



Welcome to CS162

David E. Culler

CS162 – Operating Systems and Systems Programming

<http://cs162.eecs.berkeley.edu/>

Lecture 1

August 29, 2019

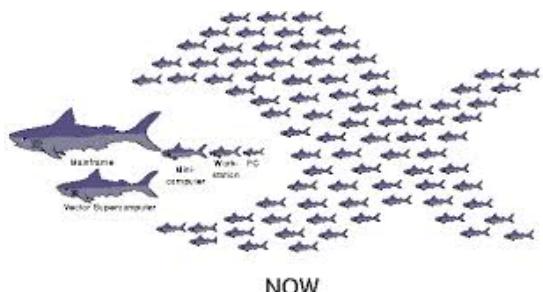
Overflow Lecture 8/30 @ 2:00 in 306 Soda

Read A&D Ch1
HW0 out, due 9/6



Introducing the CS162 Team - me

- **David Culler (culler@berkeley.edu)**
 - 783 Soda Hall
 - <http://www.cs.berkeley.edu/~culler>
 - Office hours: M 2-3, Tu 3:30-5, W 1-2
 - Before/after class
- **Extreme Systems**
 - Cray Time Sharing System
 - OS386, OS286
 - Active Messages
 - Massive Clusters for Internet Services
 - TinyOS / Berkeley Motes (IoT)
 - BOSS – OS for the Built Environment





Your TAs / Mentors



Jack Kolb
jkolb@cs.berkeley.edu



Sam Kumar
samkumar@cs.berkeley.edu



Varsha Ramakrishnan
vio@berkeley.edu



Nicholas Riasanovsky
njriasanovsky@berkeley.edu



Annie Ro
anniero@berkeley.edu



Alan Ton
alanton@berkeley.edu



William Walker
wewalker@berkeley.edu



Sharie Wang
sharie@berkeley.edu



William Wang
hwang97@berkeley.edu



Alexander Wu
alexanderwu@berkeley.edu



Yi Zhao
yi@berkeley.edu

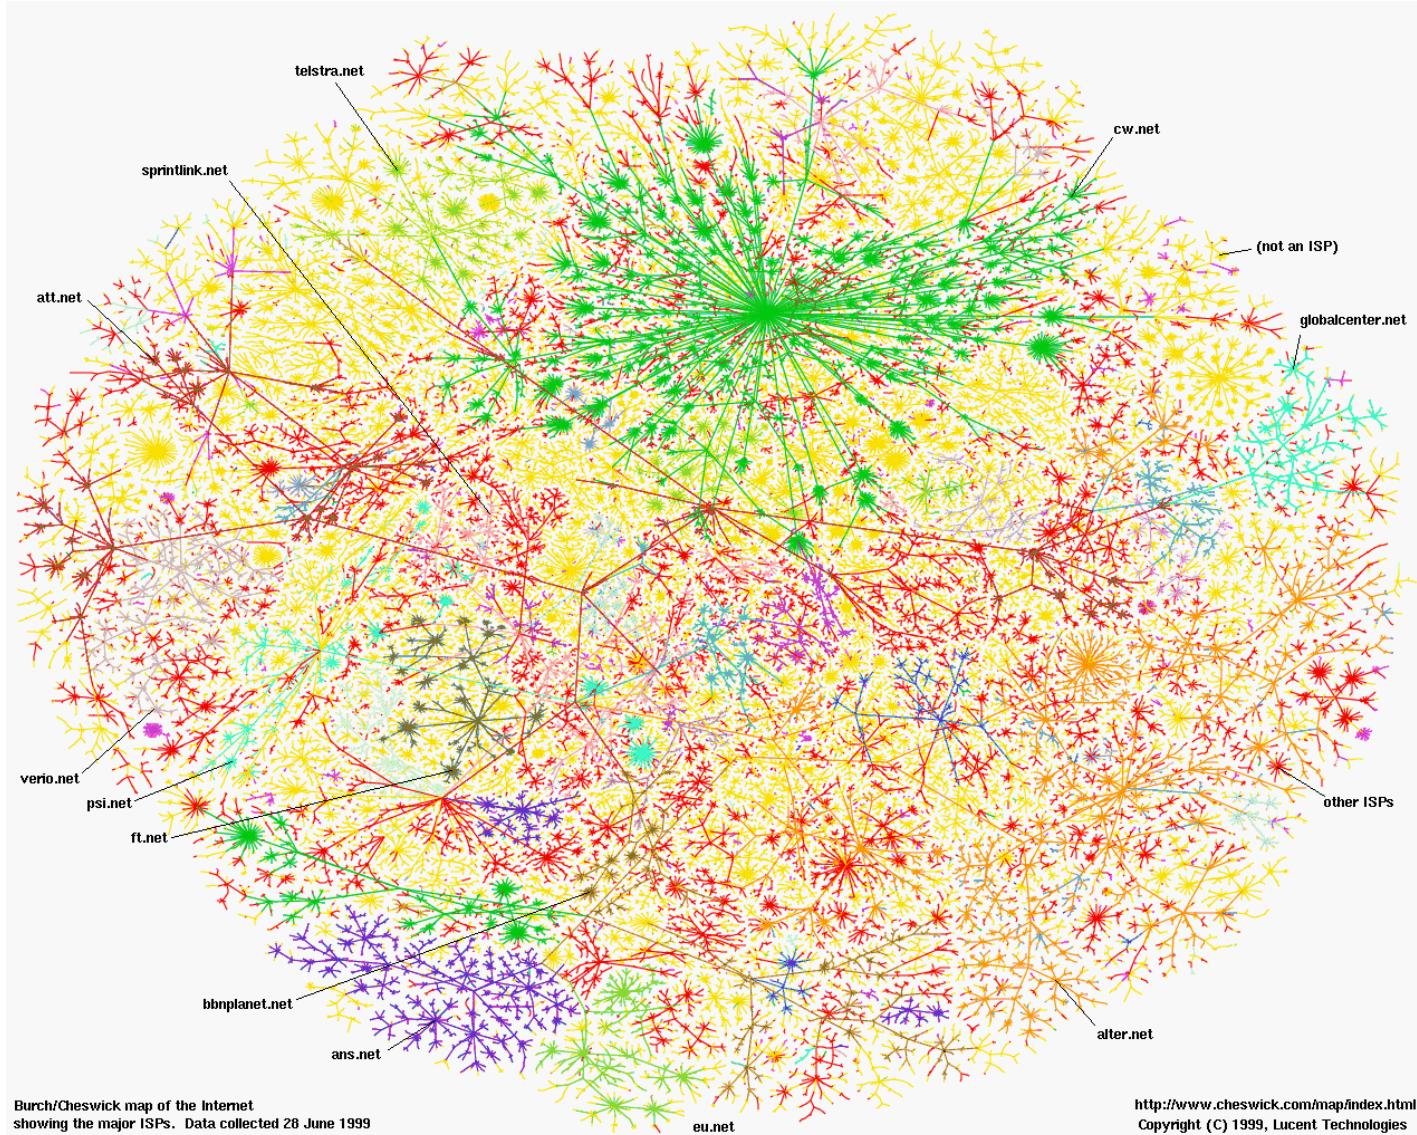


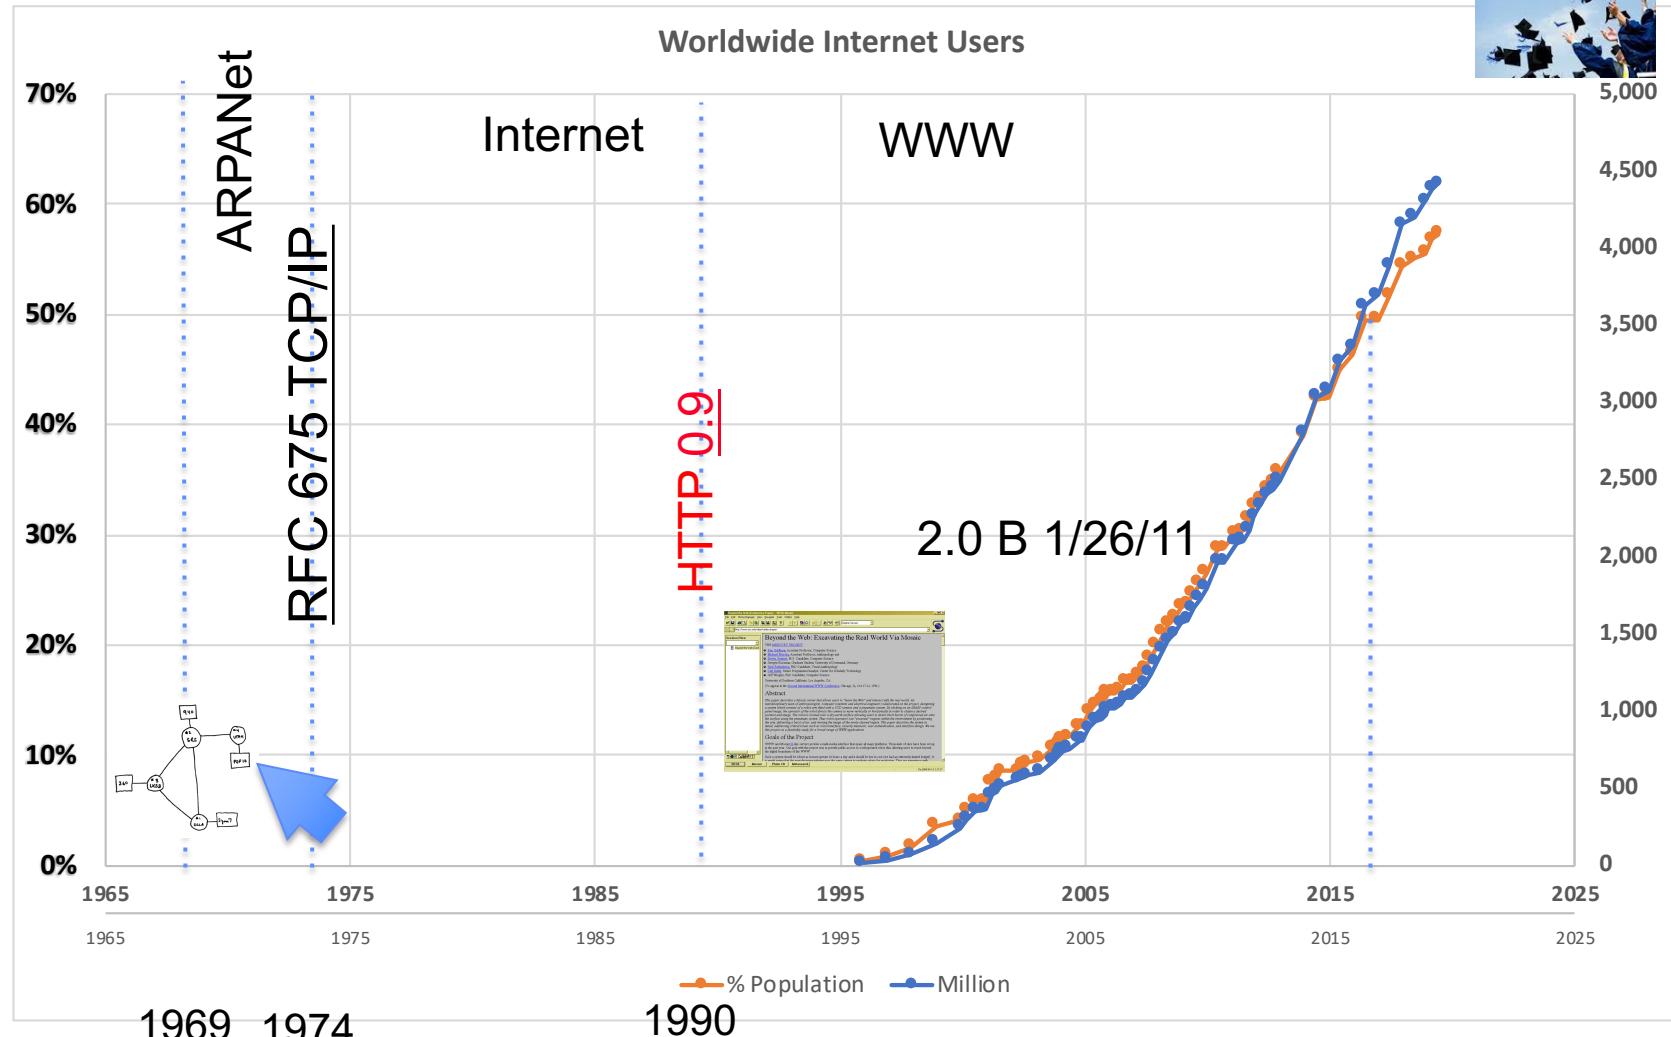
Today's Objectives

- Introduce you to Operating Systems & Their Design
- Introduce the CS162 instructional team & Plan
- Establish expectations and logistics
- Maybe get a little excited about how OS is so essential in creating and advancing this “connected world” ...



Most Transformative Artifact of Human Civilization ...



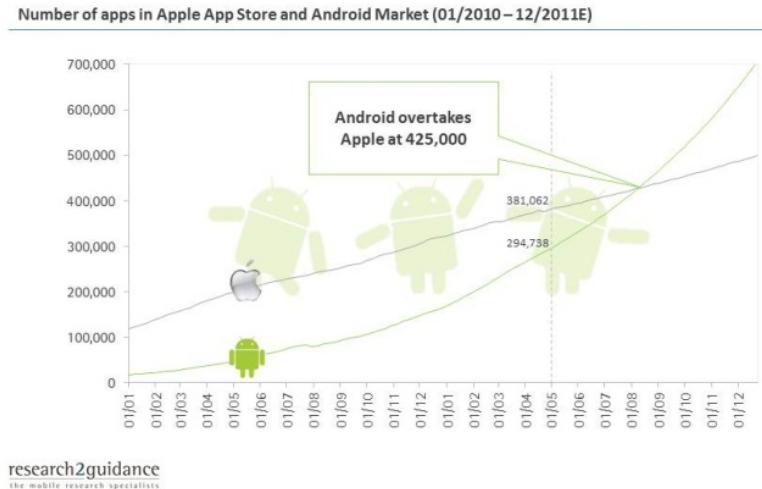




Operating Systems at the heart of it all ...

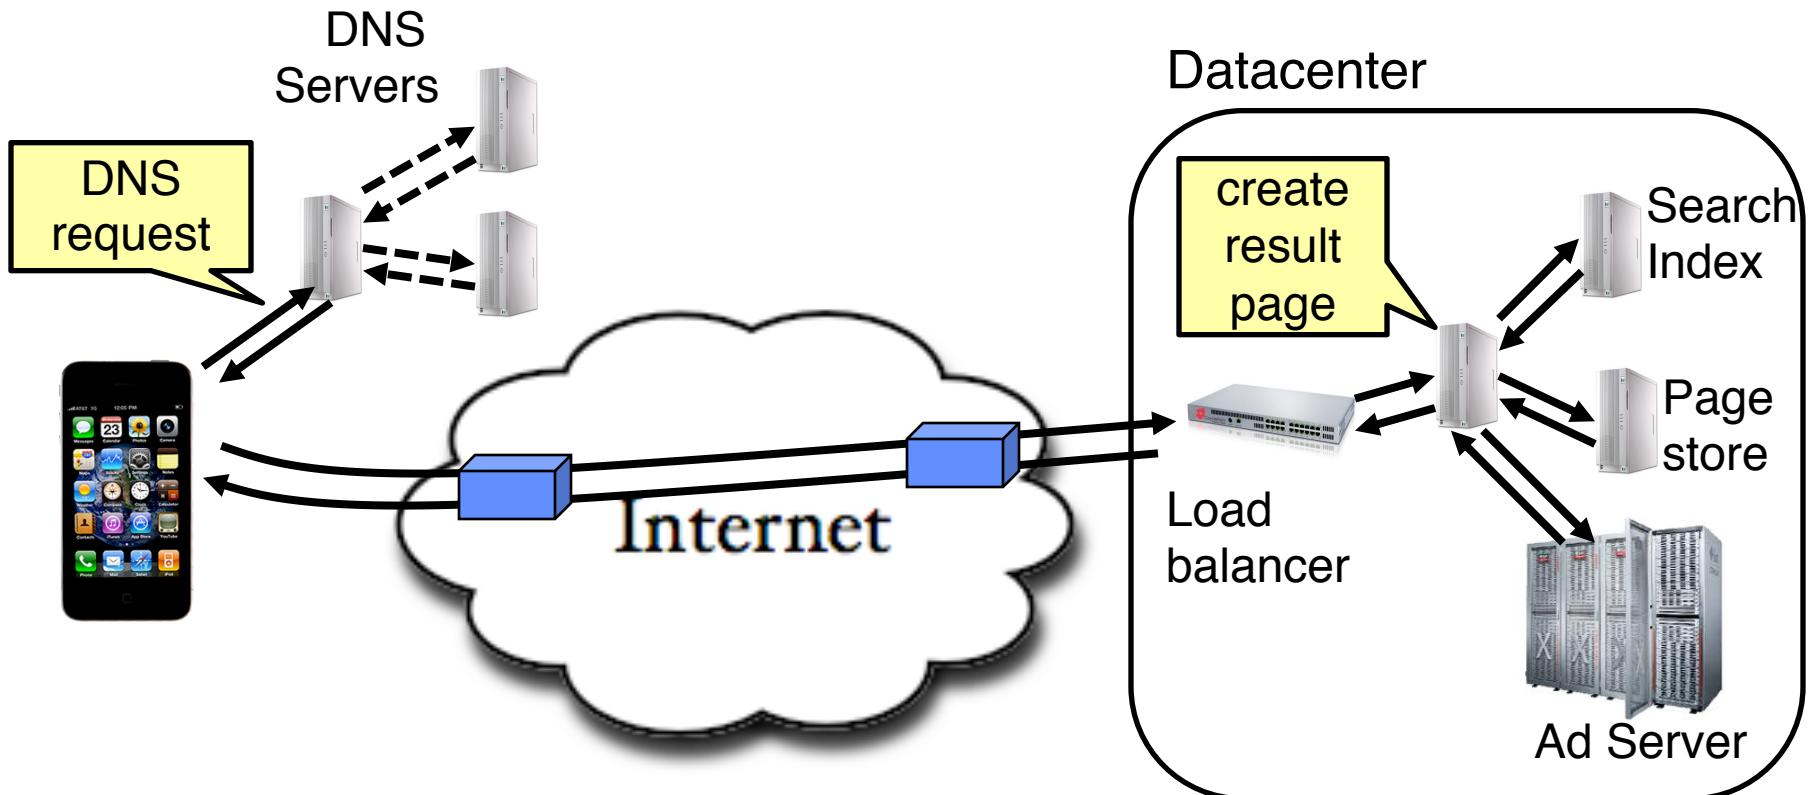
- Make the incredible advance in the underlying technology available to a rapid evolving body of applications.
 - Processing, Communications, Storage, Interaction, Protected Sharing

- The key building blocks
 - Processes, Scheduling
 - Concurrency, Coordination
 - Address spaces, Translation
 - Protection, Isolation, Sharing, Security
 - Communication, Protocols
 - Persistent storage, transactions, consistency, resilience
 - Interfaces to all devices





Example: What's in a Search Query?



- **Complex interaction of multiple components in multiple administrative domains**
 - Systems, services, protocols, ...



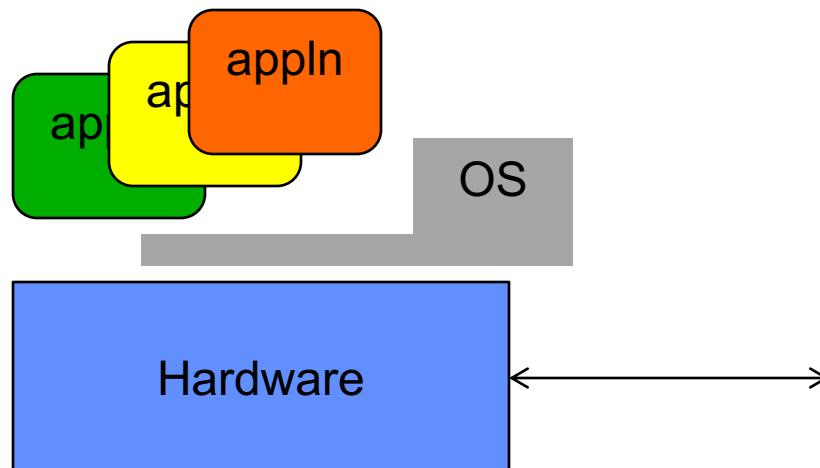
Why take CS162?

- **Some of you will may design and build operating systems or components of them.**
 - Perhaps more now than ever
- **Many of you will create systems that utilize the core concepts in operating systems.**
 - Whether you build software or hardware
 - The concepts and design patterns appear at many levels
- **All of you will build applications, etc. that utilize operating systems**
 - The better you understand their design and implementation, the better use you'll make of them.



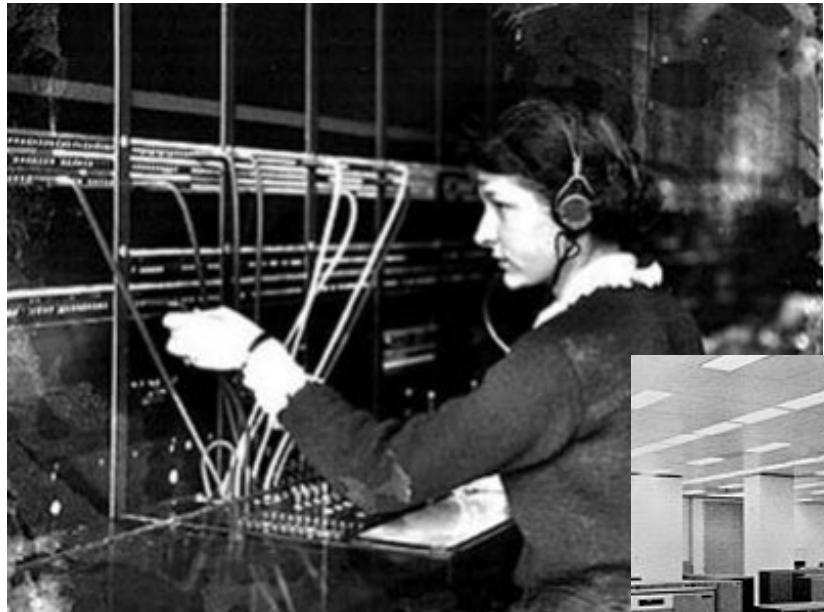
What is an operating system?

- **Special layer of software that provides application software access to hardware resources**
 - Convenient abstraction of complex hardware devices
 - Protected access to Shared resources
 - Security and Authentication
 - Communication amongst logical entities





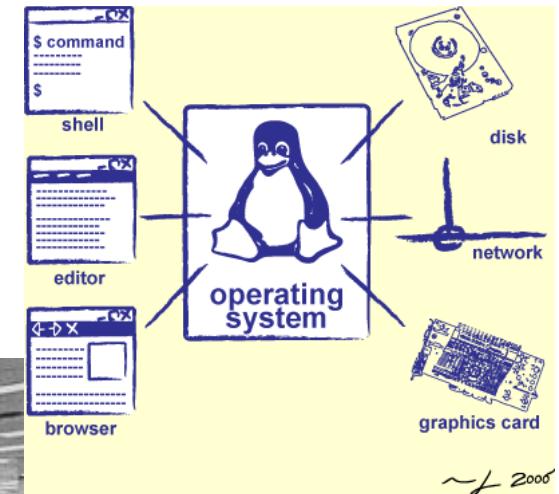
Operator ...



Switchboard Operator



Computer Operators

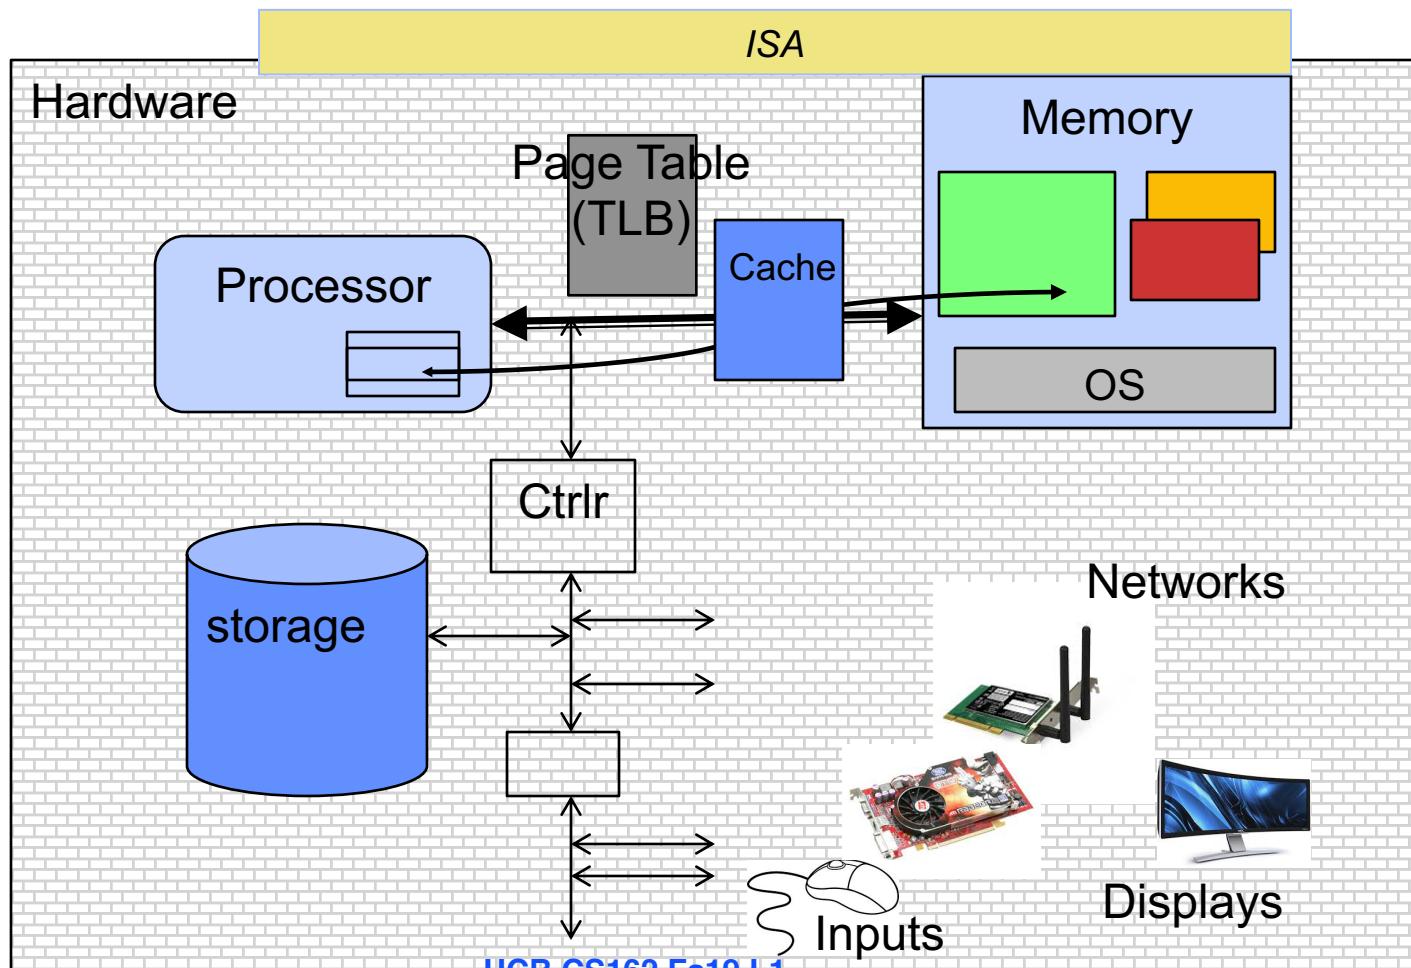




What make something a system?

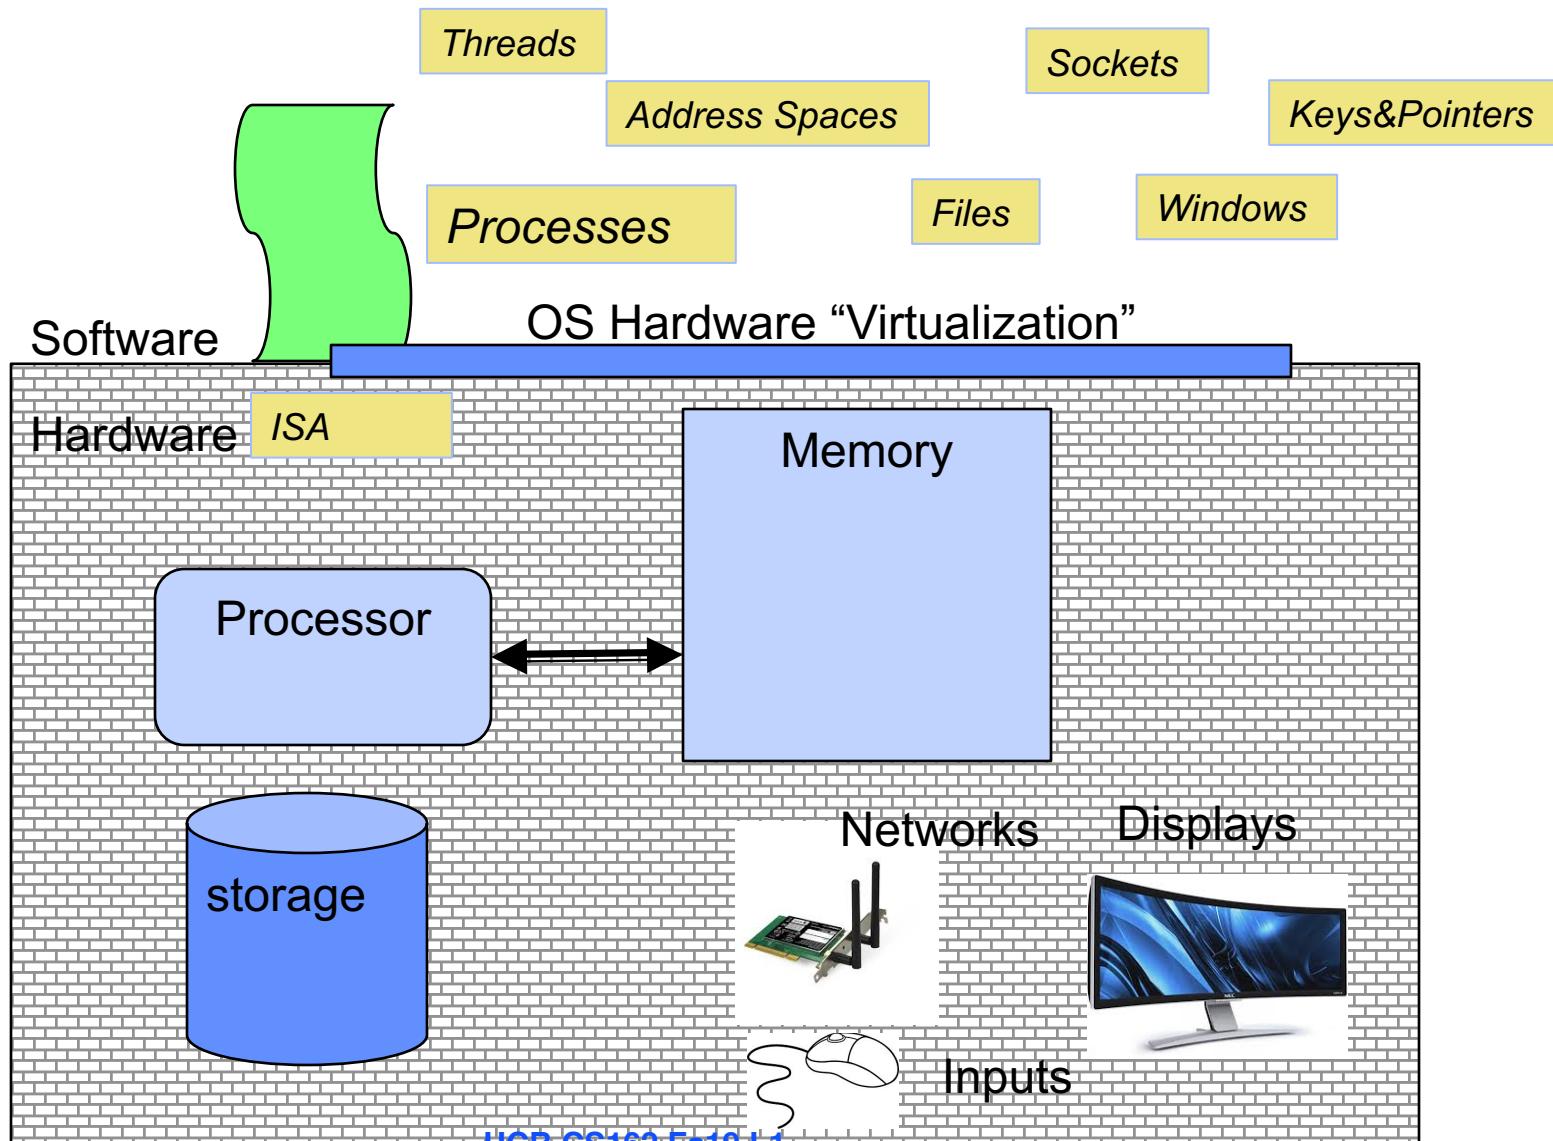
- **Multiple interrelated parts**
 - Each potentially interacts with the others, ...
- **Frame of mind:**
 - meticulous error handling, anticipating all possible failure cases, defending against malicious/careless users, ...
- **An important part of this class**

CS61C – Machine Structures (and C)





OS Basics: “Virtualizing the Machine”





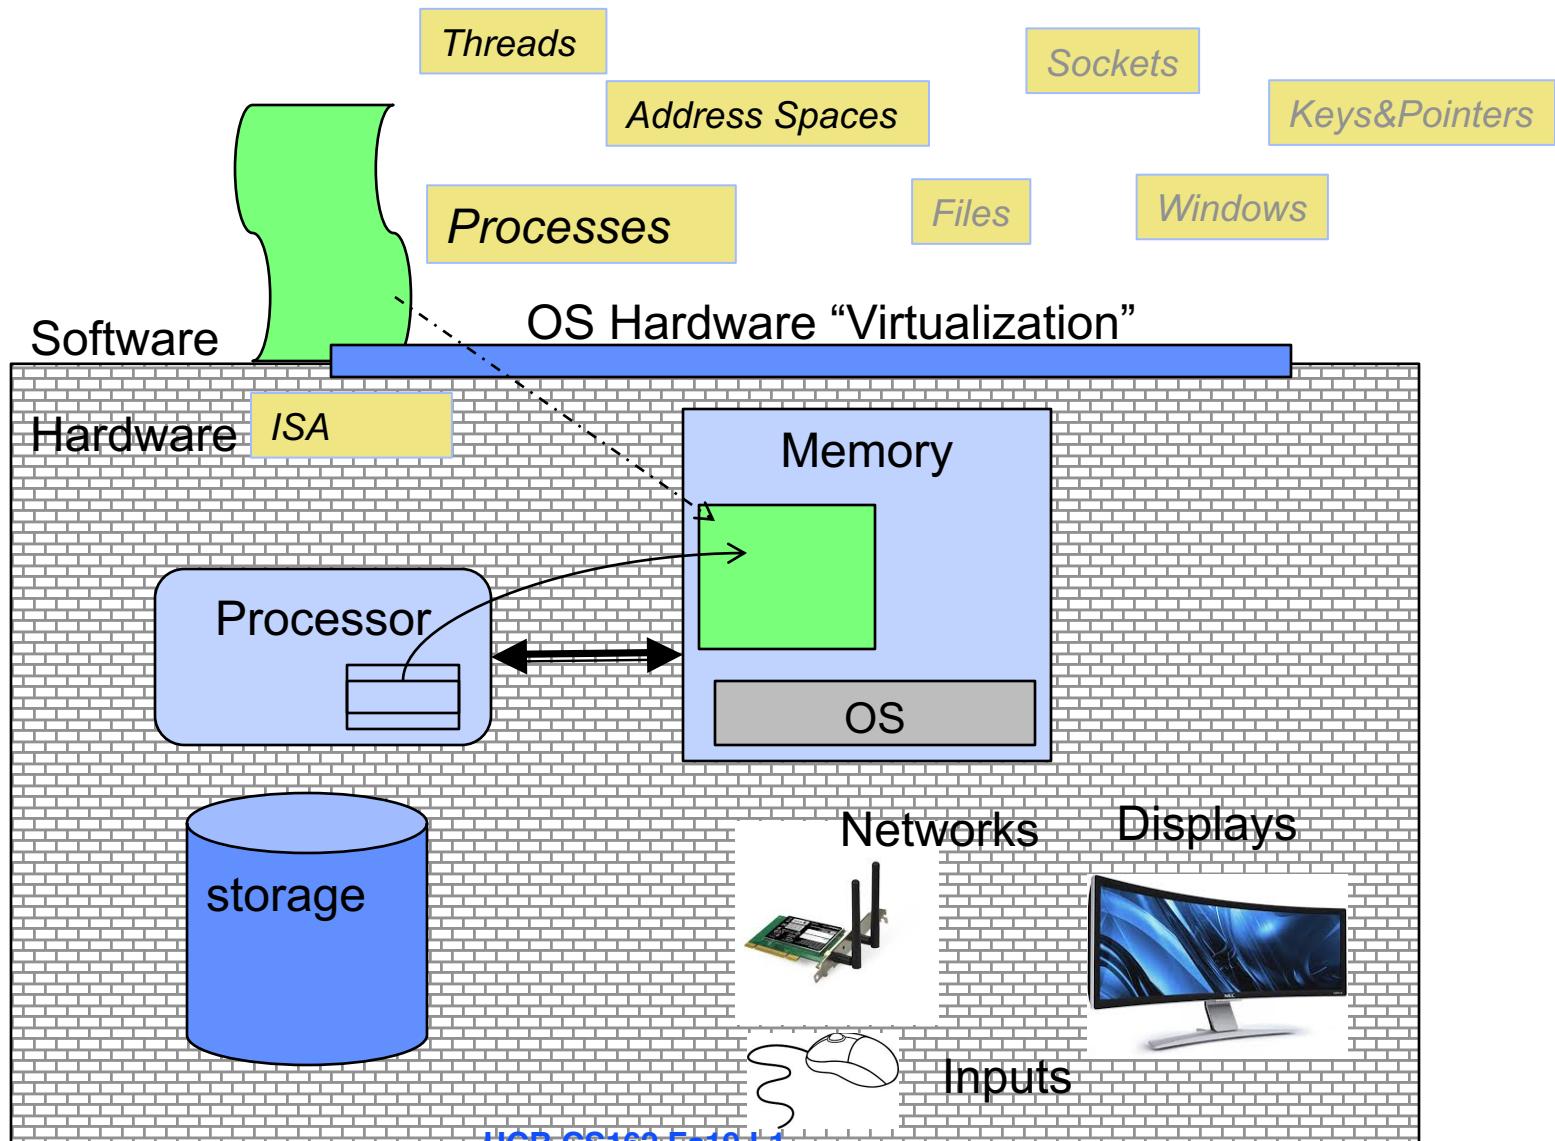
What is an Operating System?

- Illusionist
 - Provide clean, easy to use abstractions of physical resources
 - » Infinite memory, dedicated machine
 - » Higher level objects: files, users, messages
 - » Masking limitations, virtualization





OS Basics: Program => Process



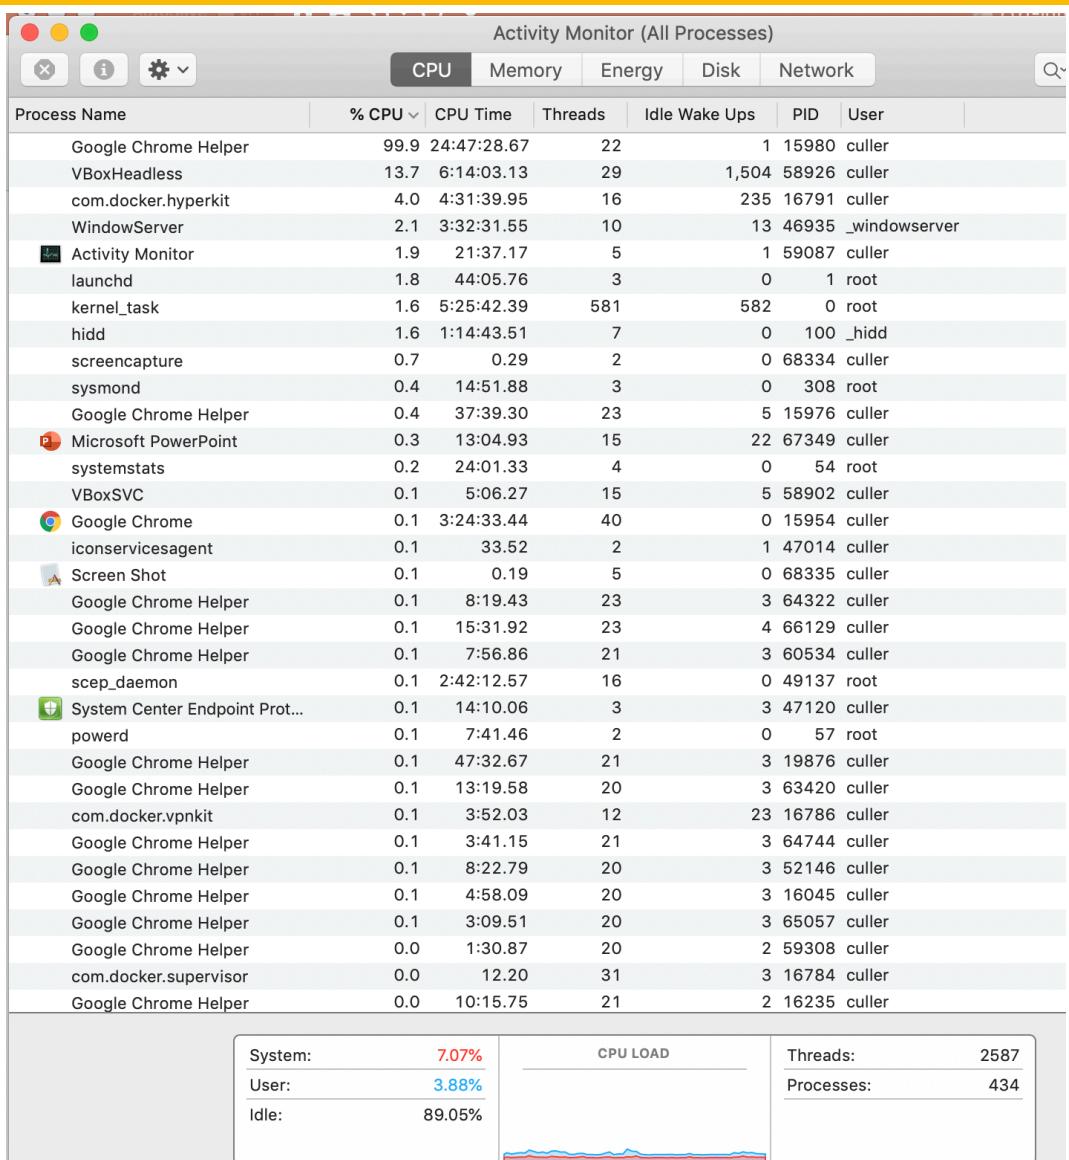


Defn: Process

- **Address Space**
- **One or more *threads* of control**
- **Additional system state associated with it**
- **Thread:**
 - locus of control (PC)
 - Its registers (processor state when running)
 - And its “stack” (SP)
 - » As required by programming language runtime

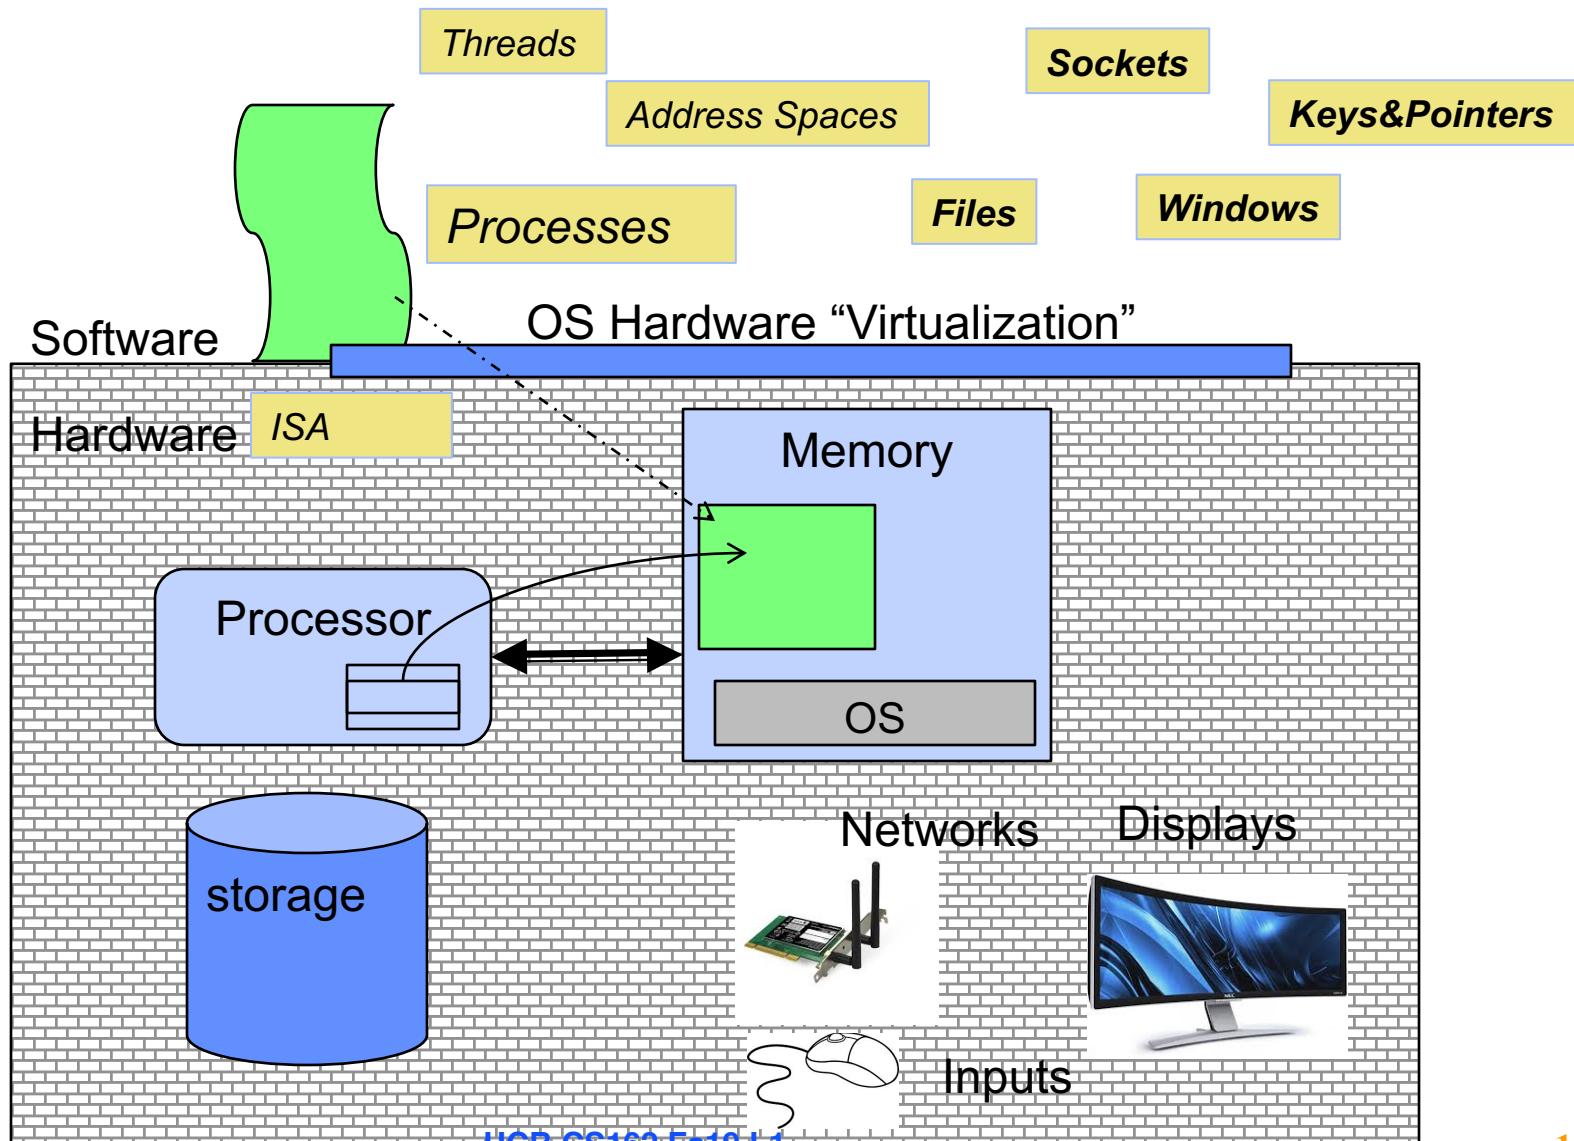


For Example ...



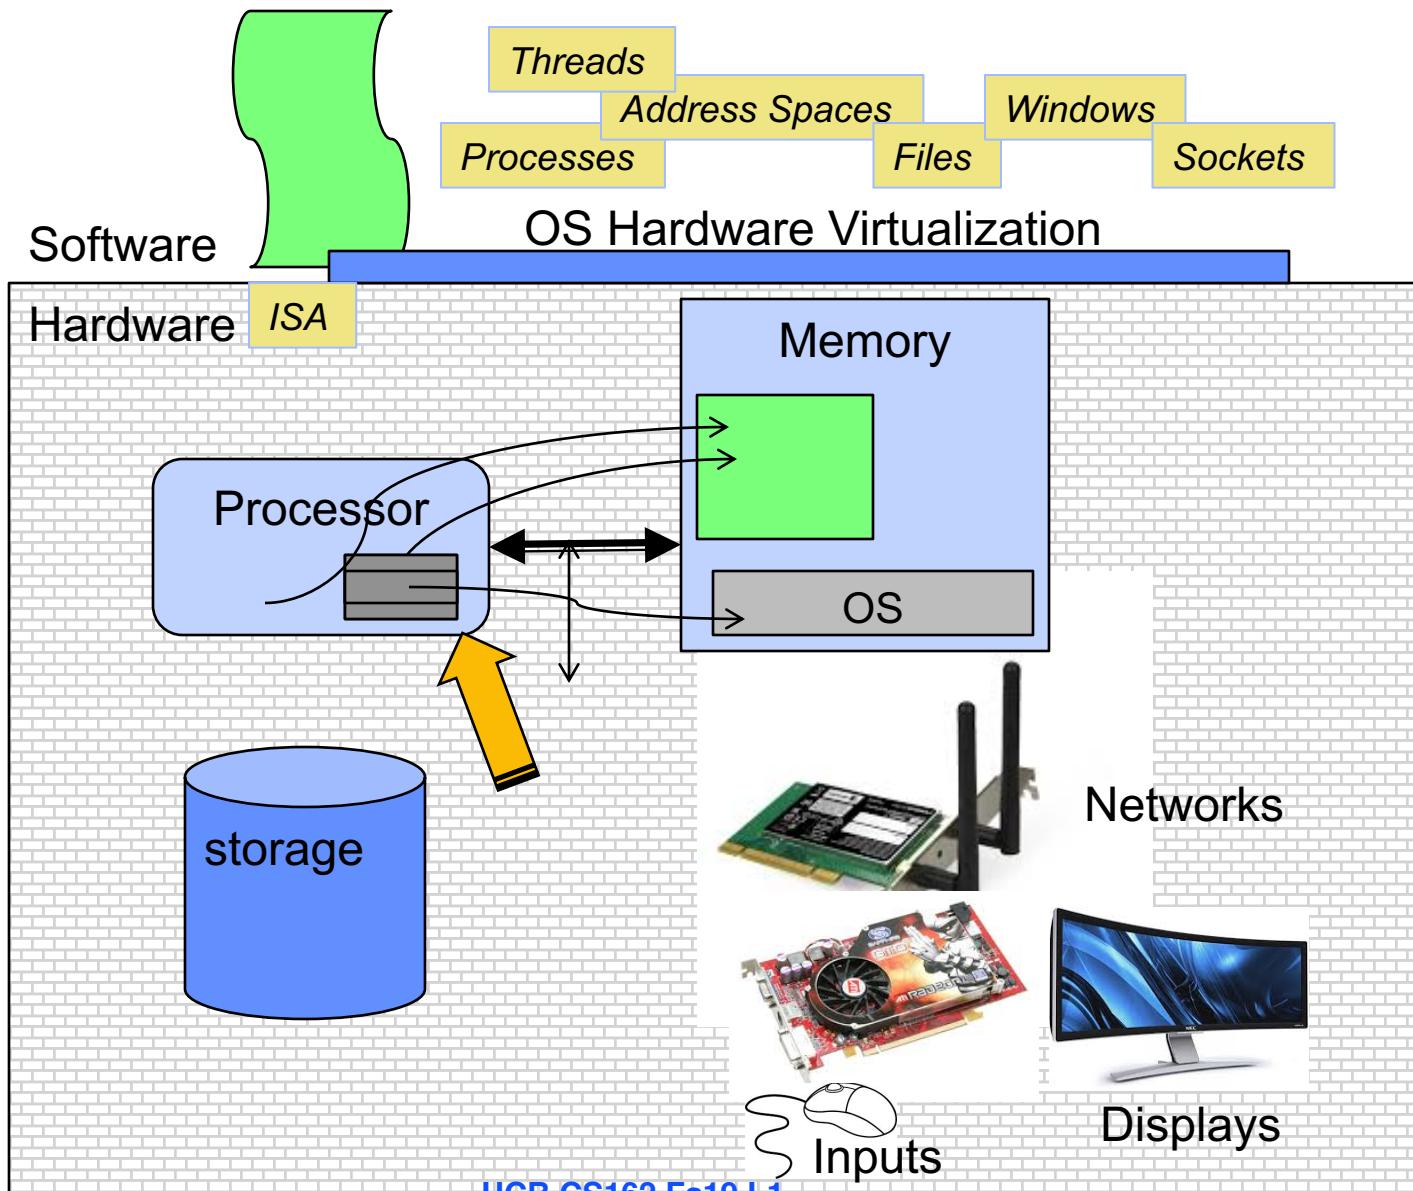


OS Basics: Program => Process



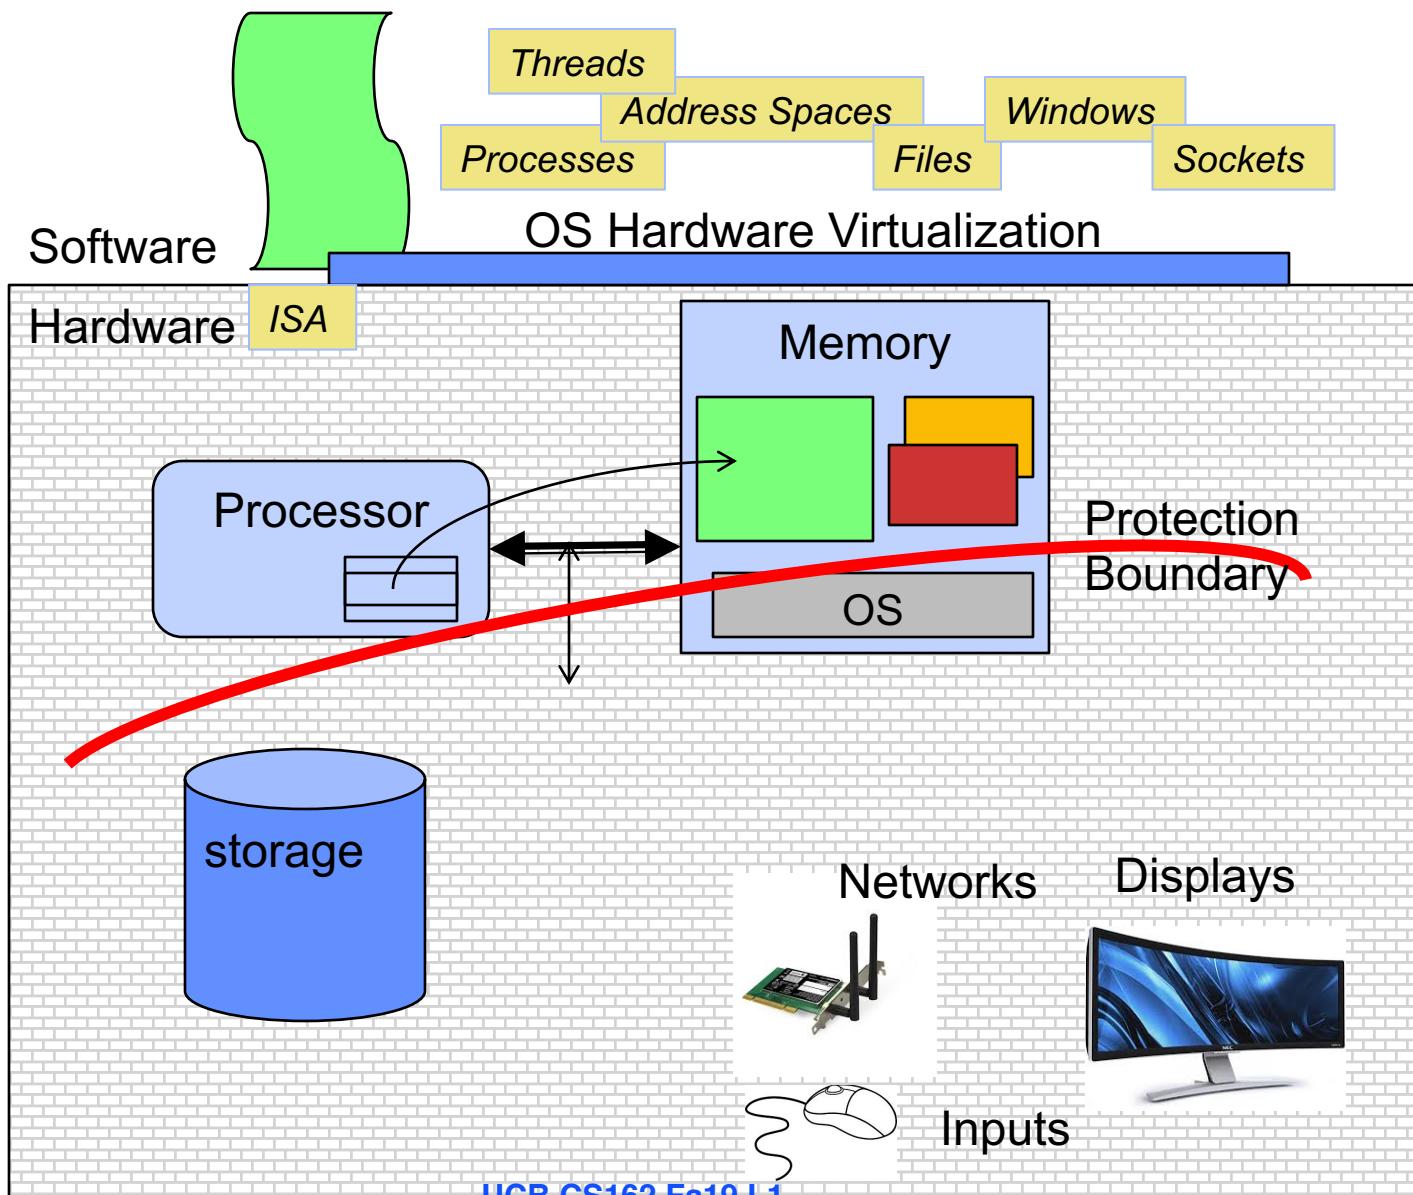


OS Basics: Context Switch





OS Basics: Scheduling, Protection





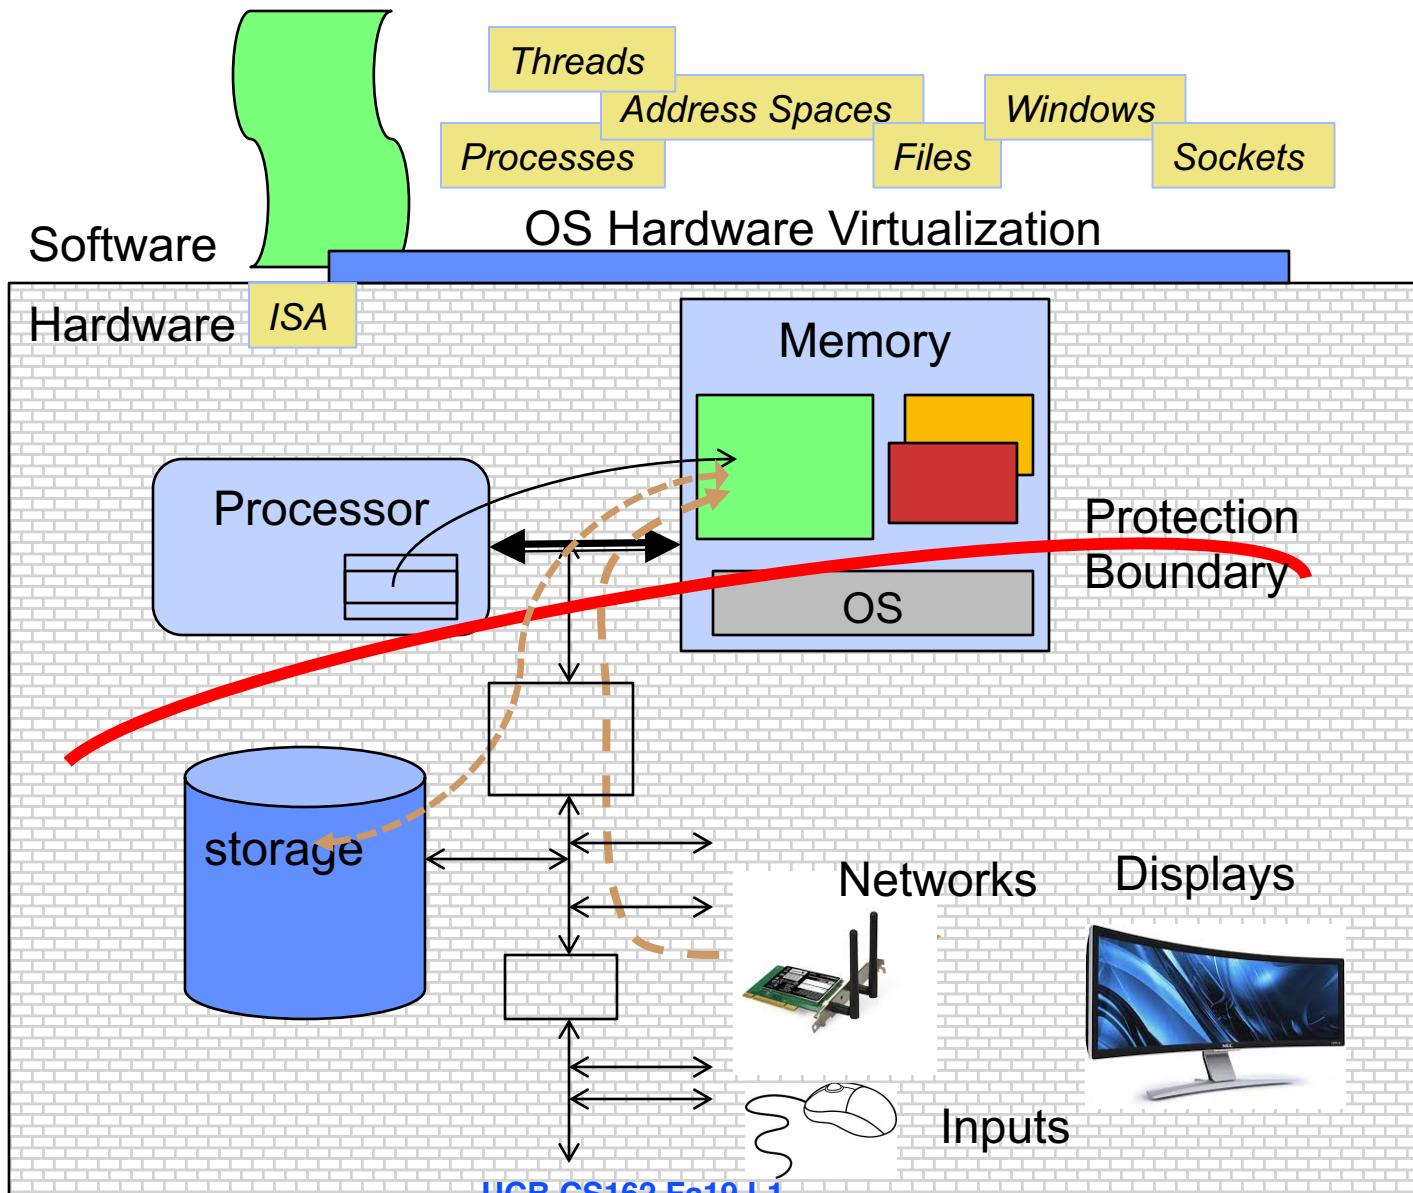
What is an Operating System?



- **Referee**
 - Manage sharing of resources, Protection, Isolation
 - » Resource allocation, isolation, communication
- **Illusionist**
 - Provide clean, easy to use abstractions of physical resources
 - » Infinite memory, dedicated machine
 - » Higher level objects: files, users, messages
 - » Masking limitations, virtualization



OS Basics: I/O





What is an Operating System?

- **Referee**

- Manage sharing of resources, Protection, Isolation
 - » Resource allocation, isolation, communication

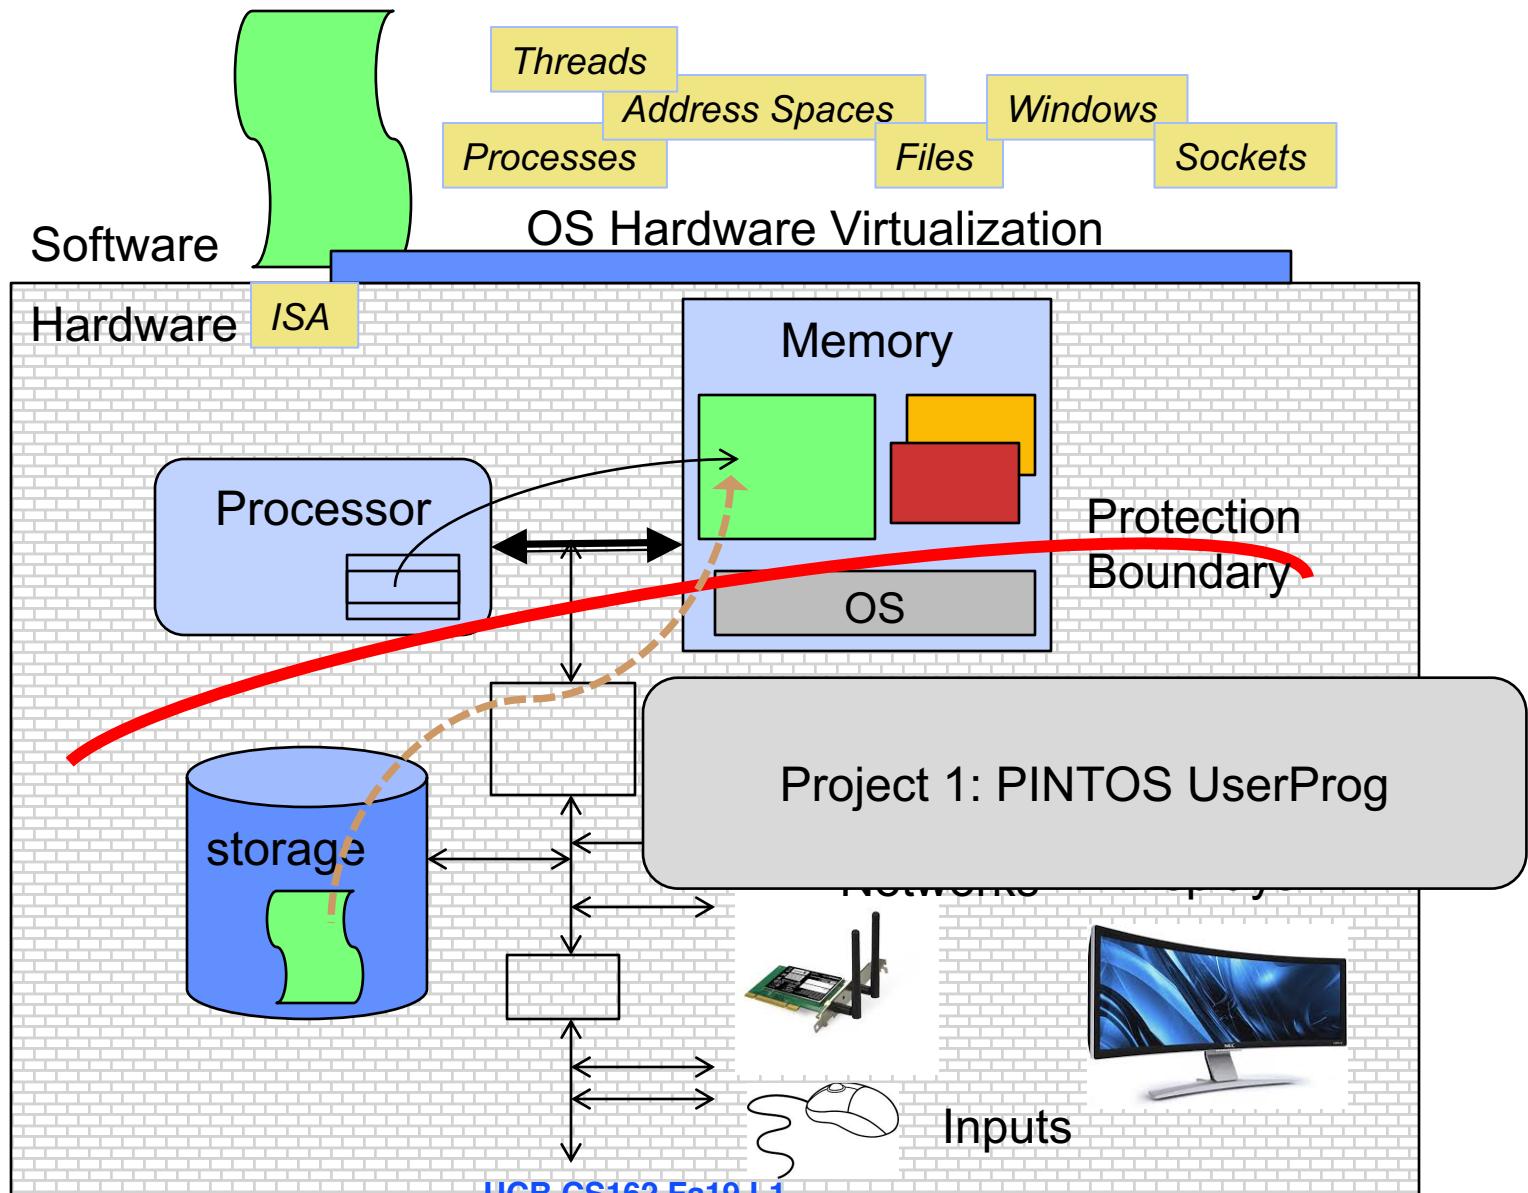
- **Illusionist**

- Provide clean, easy to use abstractions of physical resources
 - » Infinite memory, dedicated machine
 - » Higher level objects: files, users, messages
 - » Masking limitations, virtualization

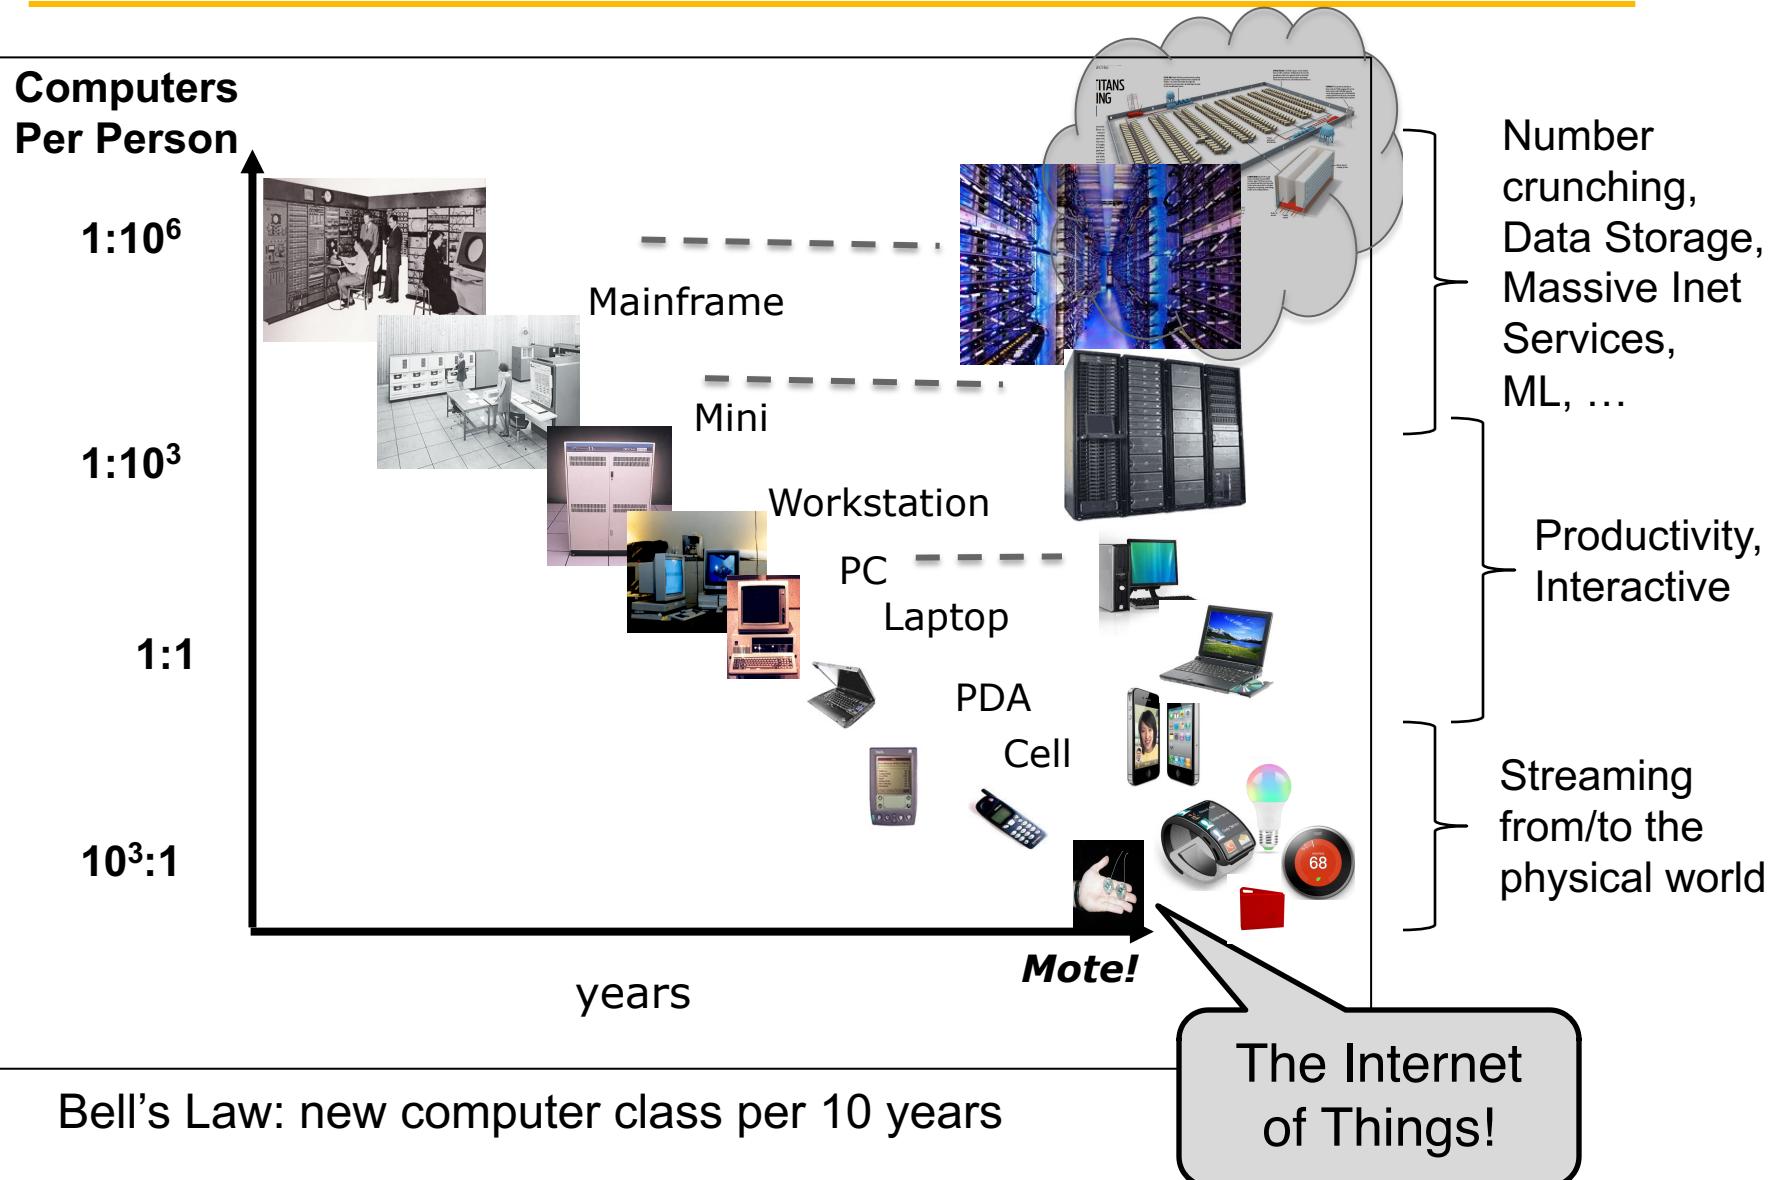
- **Glue**

- Common services
 - » Storage, Window system, Networking
 - » Sharing, Authorization
 - » Look and feel

Creating Process & Loading a Program



Across incredibly diversity



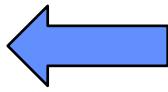


And Range of Timescales

Jeff Dean: "Numbers Everyone Should Know"

L1 cache reference	0.5 ns
Branch mispredict	5 ns
L2 cache reference	7 ns
Mutex lock/unlock	25 ns
Main memory reference	100 ns
Compress 1K bytes with Zippy	3,000 ns
Send 2K bytes over 1 Gbps network	20,000 ns
Read 1 MB sequentially from memory	250,000 ns
Round trip within same datacenter	500,000 ns
Disk seek	10,000,000 ns
Read 1 MB sequentially from disk	20,000,000 ns
Send packet CA->Netherlands->CA	150,000,000 ns

Key Stroke / Click
100 ms



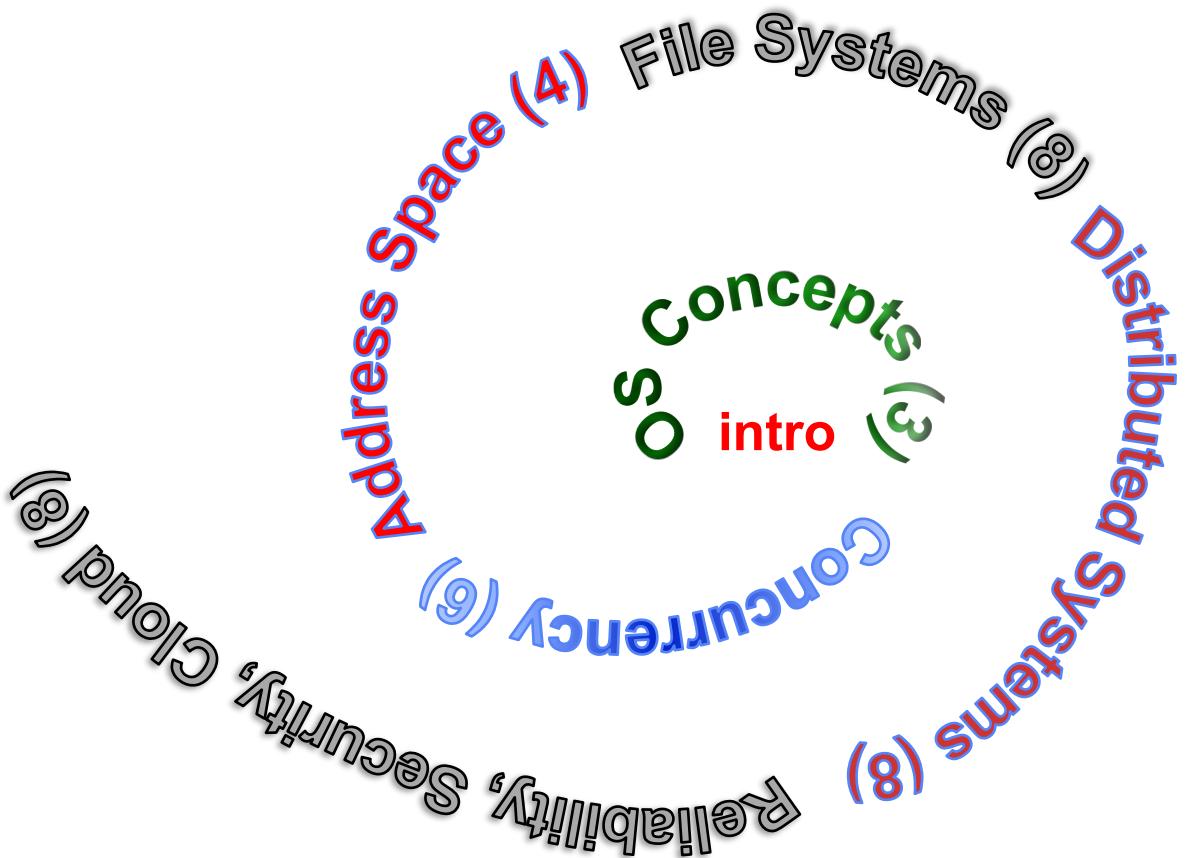


Operating systems

- Provide consistent abstractions to apps, even on diverse hardware
 - File systems, Window Systems, Communications, ...
 - Processes, threads
 - VMs, containers
 - Naming systems
- Manage resources shared among multiple applications:
 - Memory, CPU, storage, ...
- Achieved by specific algorithms and techniques:
 - Scheduling
 - Concurrency
 - Transactions
 - Security
- Across immense scale – from 1's to Billions



Course Structure: Spiral





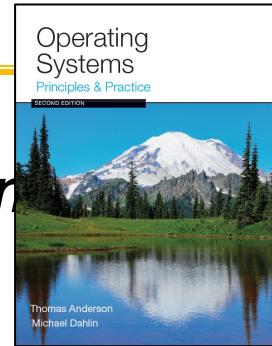
Syllabus – in nutshell

- **OS Concepts**
 - Process, I/O, Networks and VM
- **Concurrency**
 - Threads, scheduling, locks, deadlock, scalability, fairness
- **Address Space**
 - Virtual memory, address translation, protection, sharing
- **File Systems**
 - i/o devices, file objects, storage, naming, caching, performance, paging, transactions, databases
- **Distributed Systems**
 - Protocols, Tiers, RPC, NFS, DHTs, Consistency, Scalability, multicast
- **Reliability & Security**
 - Fault tolerance, protection, security
- **Cloud Infrastructure**



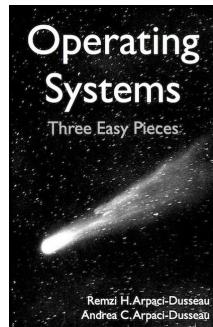
Textbook & Resources

- Course text: *Operating Systems, Principles and Practice*, 2nd ed., by Anderson and Dahlin

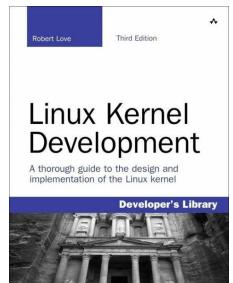


- Recommended Supplementary Material:

- *Operating Systems: Three Easy Pieces*, by Remzi and Andrea Arpaci-Dusseau, available for free online at:
<http://pages.cs.wisc.edu/~remzi/OSTEP>
 - *Linux Kernel Development*, 3rd ed., by Robert Love



- Additional materials on course website "Readings" page





Infrastructure

- Course Website:
<https://cs162.eecs.berkeley.edu>
- Piazza for Q&A:
<https://piazza.com/berkeley/spring2019/cs162>
- Anonymous Feedback Form:
 - <http://bit.ly/cs162fa19anon>



Learning by Doing

- **Three Group Projects (4 weeks, Pintos in C)**
 - 1. User-programs (exec & syscall)
 - 1. Threads & Scheduling
 - 3. File Systems
- **Individual Homeworks (2 weeks) - preliminary**
 - 0. Tools & Environment, Autograding, recall C, executable
 - 1. Getting to Pintos-scale C, Basic Threads
 - 2. BYOS – build your own shell
 - 3. Threads and Synchronization
 - 4. Sockets & Threads in HTTP server
 - 5. Memory mapping and management
 - 6. Distributed Computing – KV store in Kubernetes
 - 7. Cloud APIs, GRPC
- **Weekly quick surveys**
 - Ungraded, reinforce material, instructional diagnostic



Steeping in GO ...

- Last one or two homeworks in Go
- Nudges to learn Go at your own pace along the way



Group Projects

- **Project teams have 4 members**
 - never 5, 3 req's serious justification)
 - Must work in groups in “the real world”
 - Same section (at least same TA)
- **Communication and cooperation will be essential**
 - Design Documents
 - Slack/Messenger/whatever doesn't replace face-to-face
- **Everyone should do work and have clear responsibilities**
 - You will evaluate your teammates at the end of each project
 - Dividing up by Task is the worst approach. Work as a team.
- **Communicate with supervisor (TAs)**
 - What is the team's plan?
 - What is each member's responsibility?
 - Short progress reports are required



Grading

- **36% exams**
- **36% projects**
- **22% homework**
- **6% participation**
- **Project grading**
 - Initial design document, Design review, Code, Final design
- **Submission via *git push* to release branch**
 - Triggers autograder
- **Regular *git push* so TA sees your progress**



Getting started

- Start homework 0 immediately
 - Get **cs162-xx** account via <https://inst.eecs.berkeley.edu/webacct>
 - Github account
 - Registration survey
 - Vagrant virtualbox – VM environment for the course
 - » Consistent, managed environment on your machine
 - Get familiar with all the cs162 tools
 - Submit to autograder via git
- Go to **any** section this week
- Waitlist
 - Pull hw0, get inst acct, go to a section
 - Only enrolled students will form project groups
 - If cs162 is not for you now, please allow others to take it
 - We are going to try to grow more to enroll majors from current waitlist



Preparing Yourself for this Class

- The projects will require you to be very comfortable with programming and debugging C
 - Pointers (including function pointers, `void*`)
 - Memory Management (`malloc`, `free`, stack vs heap)
 - Debugging with GDB
- You will be working on a larger, more sophisticated code base than anything you've likely seen in 61C
- C / 61C review session Tuesday 9/3 7-8pm @ 310 Jacobs

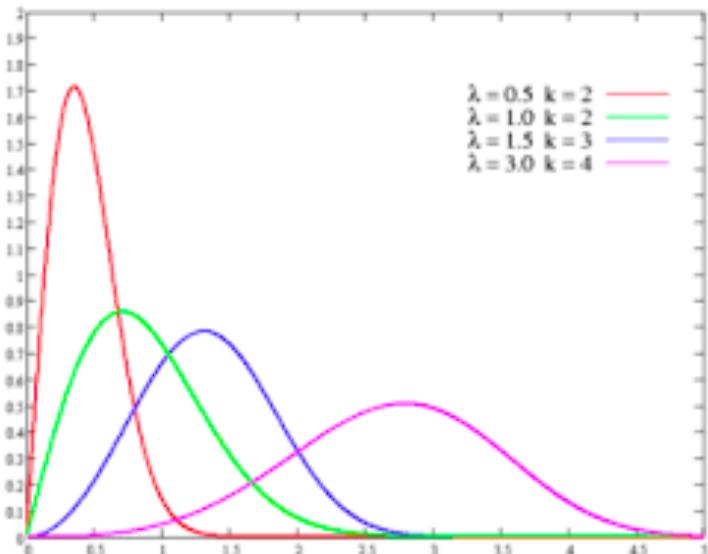


Preparing Yourself for this Class

- "Resources" page on course website
- C programming reference (still in beta):
 - <https://cs162.eecs.berkeley.edu/ladder/>
- This week's section dedicated to programming and debugging review
 - Attend ANY section this week
- Review session Tuesday evening
 - 306 Soda



Personal Growth ...



- **Homeworks will try to offer optional “warm up” and “stretch” problems so you can tailor your learning to your situation**



Personal Integrity

- **UCB Academic Honor Code:** "As a member of the UC Berkeley community, I act with honesty, integrity, and respect for others."

<http://asuc.org/honorcode/resources/HC%20Guide%20for%20Syllabi.pdf>



Collaboration Policy

OK:

**Explaining a concept
to someone in another
group**

**Discussing
algorithms/testing
strategies with other
groups**

**Helping debug
someone else's code
(in another group)**

**Searching online for
generic algorithms
(e.g., hash table)**

Not OK:

**Sharing code or test
cases with another
group**

**Copying OR reading
another group's code
or test cases**

**Copying OR reading
online code or test
cases from prior years**



Late Policy

- Deadlines are at 9:00 PM Pacific Time
- 4 slip days to use for homeworks, no more than 1 per assignment
- 4 slip days to use (as a group) for projects code, no more than 2 per project (no slip on report)
- No credit for late work beyond these



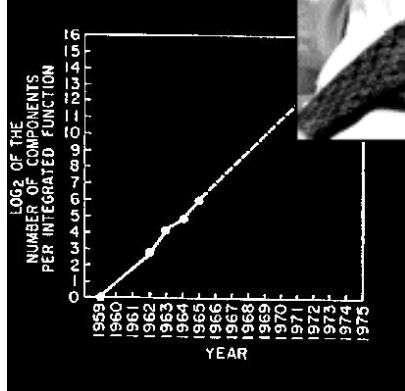
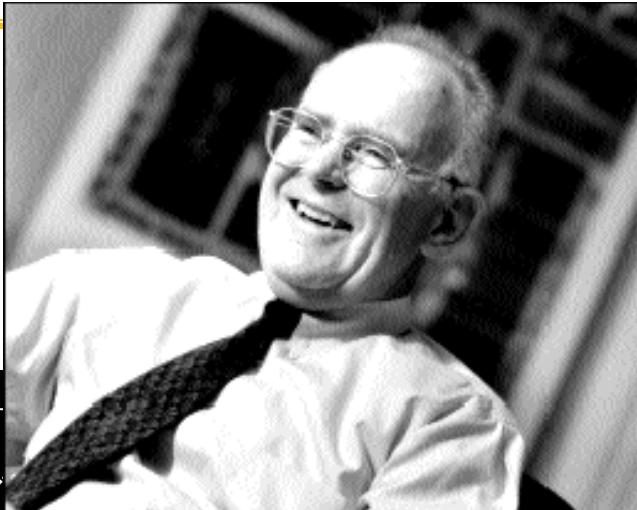
Questions



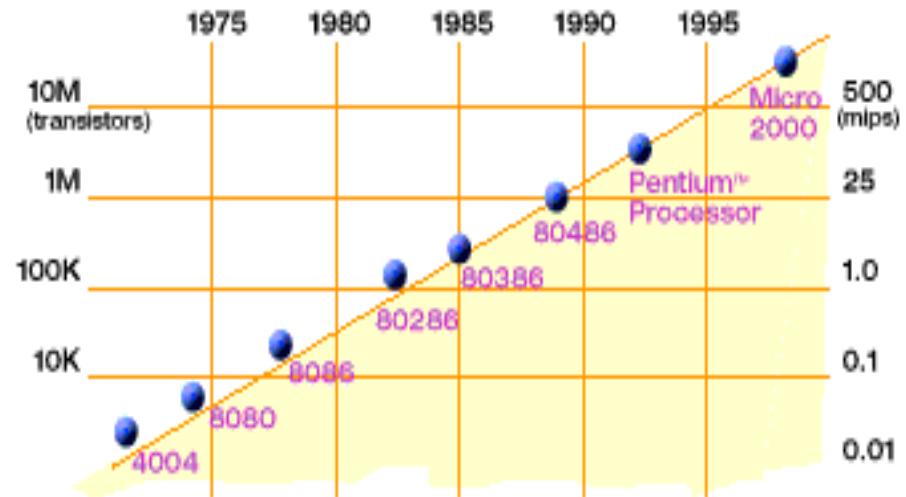
What make Operating Systems So Exciting and Challenging ?

Managing complexity

Technology Trends: Moore's Law



Gordon Moore (co-founder of Intel) predicted in 1965 that the transistor density of semiconductor chips would double roughly every 18 months.

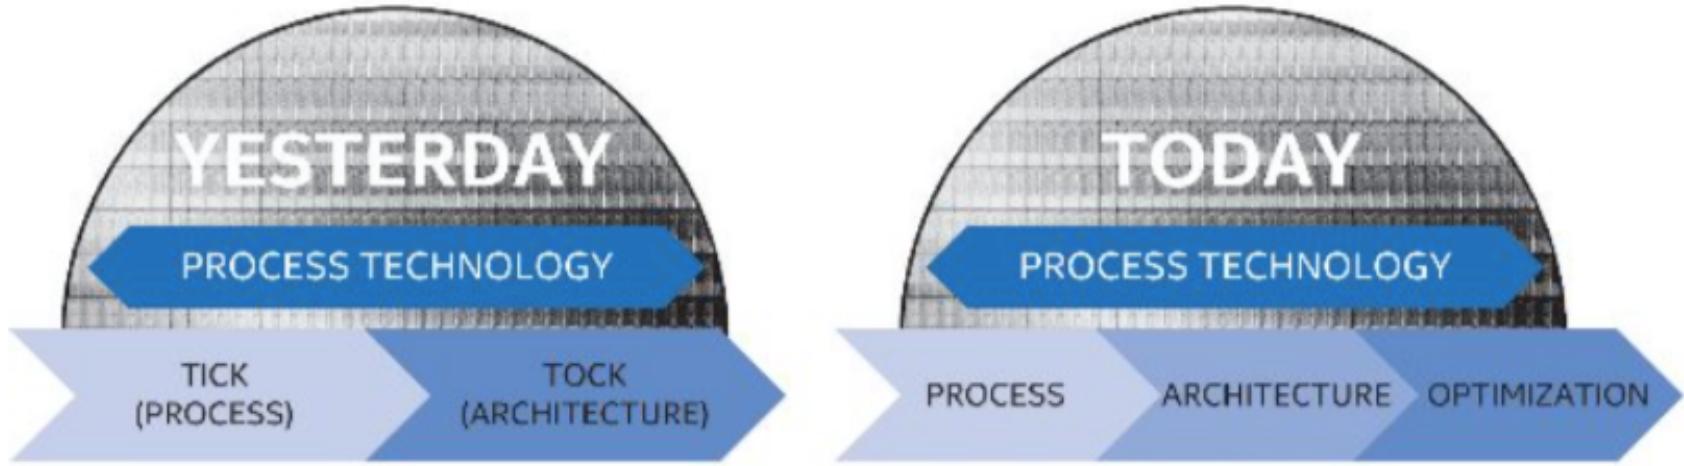


2X transistors/Chip Every 1.5 years
Called "Moore's Law"

Microprocessors have become smaller, denser, and more powerful.



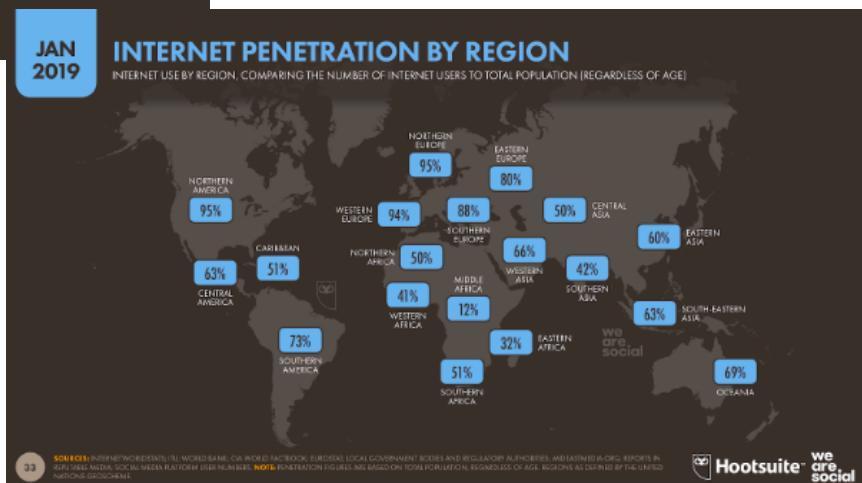
But then Moore's Law Ended...



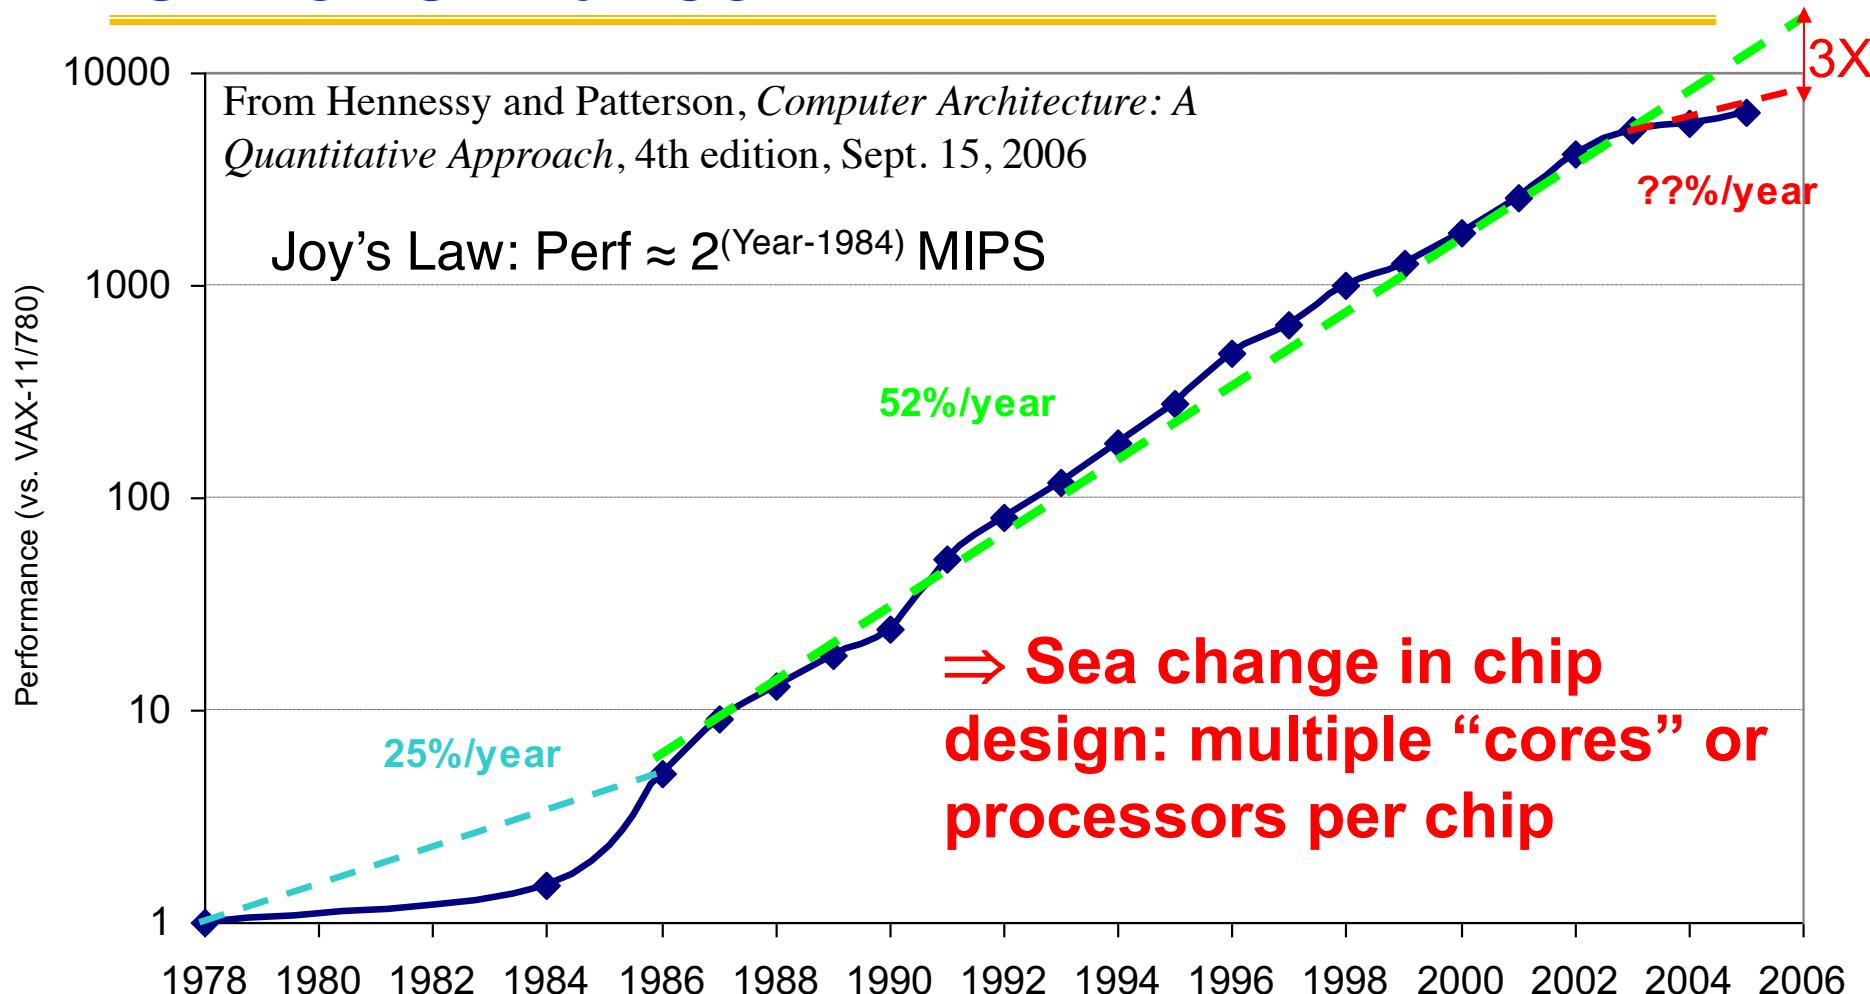
- **Moore's Law has (officially) ended -- Feb 2016**
 - No longer getting 2 x transistors/chip every 18 months...
 - or even every 24 months
- **May have only 2-3 smallest geometry fabrication plants left:**
 - Intel and Samsung and/or TSMC
- **Vendors moving to 3D stacked chips**
 - More layers in old geometries



Society connected

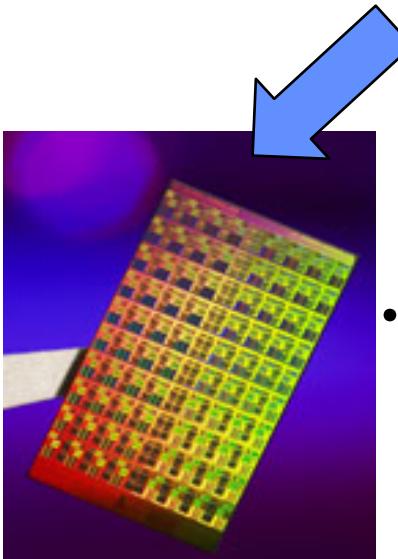


New Challenge: Slowdown in Joy's law of Performance



- VAX : 25%/year 1978 to 1986
- RISC + x86: 52%/year 1986 to 2002
- RISC + x86: ??%/year 2002 to present

ManyCore Chips: The future is here



- Intel 80-core multicore chip (Feb 2007)

- 80 simple cores
- Two FP-engines / core
- Mesh-like network
- 100 million transistors

- Intel Single-Chip Cloud Computer (August 2010)

- 24 “tiles” with two cores/tile
- 24-router mesh network
- 4 DDR3 memory controllers
- Hardware support for message-passing

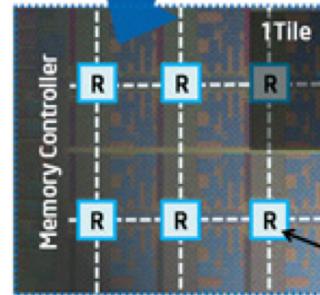
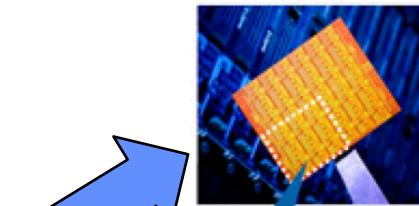
- “ManyCore” refers to many processors/chip

- 64? 128? Hard to say exact boundary

- How to program these?

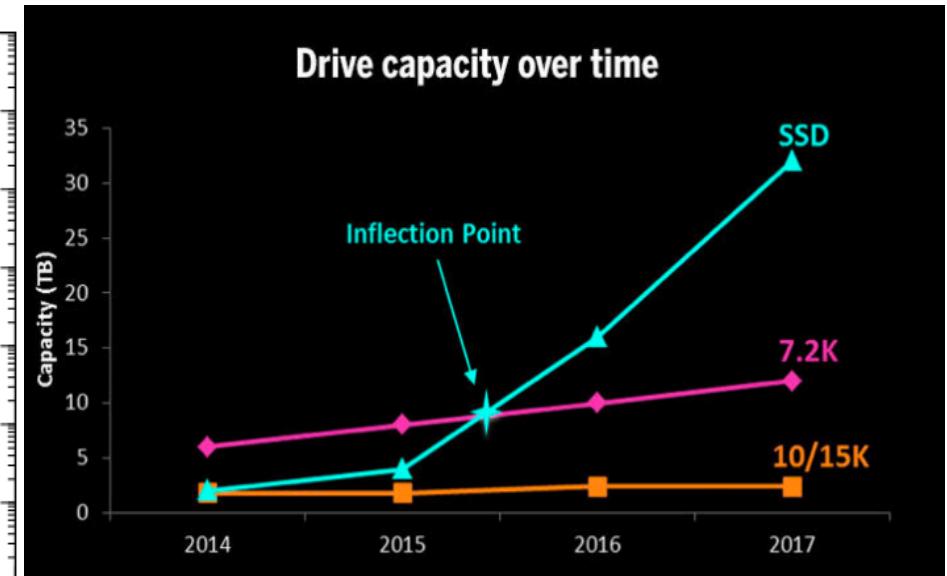
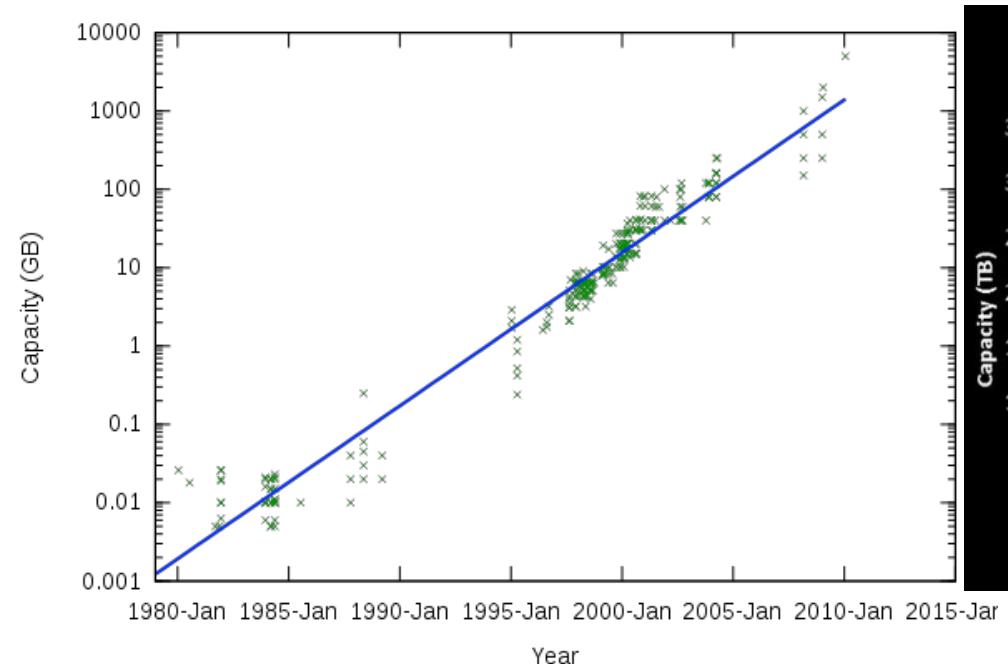
- Use 2 CPUs for video/audio
- Use 1 for word processor, 1 for browser
- 76 for virus checking???

- Parallelism must be exploited at all levels

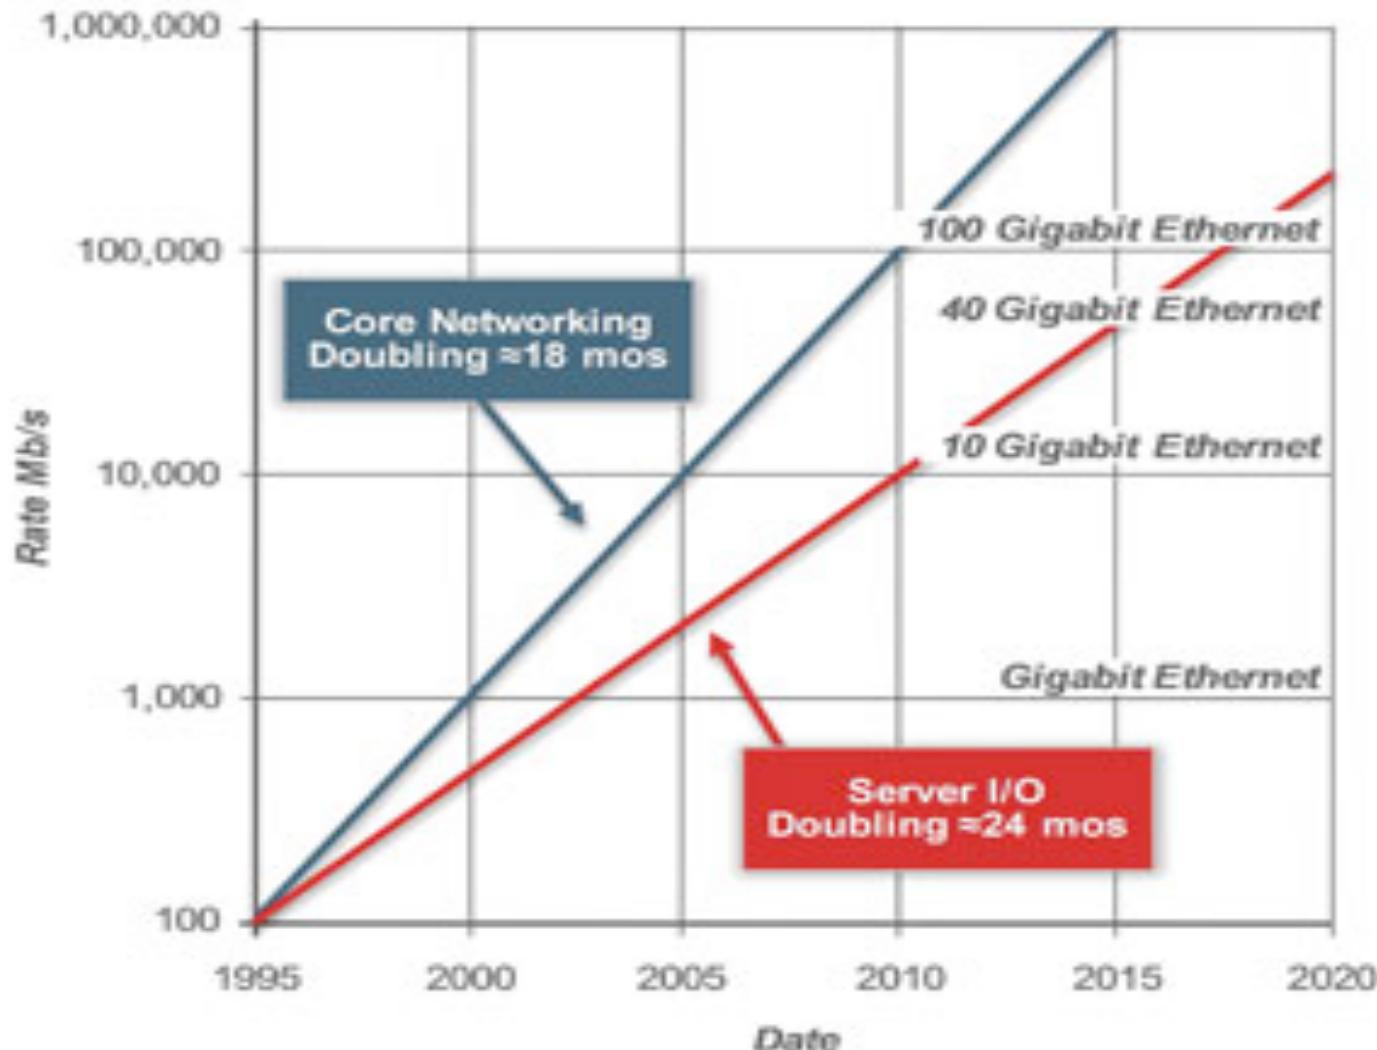




Storage Capacity



Network Capacity



(source: <http://www.ospmag.com/issue/article/Time-Is-Not-Always-On-Our-Side>)



Not Only PCs connected to the Internet

- In 2011, smartphone shipments exceeded PC shipments!

1.53B in 2017

- ~~2011 shipments:~~

– 487M smartphones



262.5M in 2017

– 414M PC clients

» 210M notebooks

164M in 2017

» 112M desktops

» 63M tablets

– 25M smart TVs

39.5M in 2017

- 4 billion phones in the world → smartphones over next few years
- Then...

Societal Scale Information Systems



- The world is a large distributed system

- Microprocessors in everything
- Vast infrastructure behind them

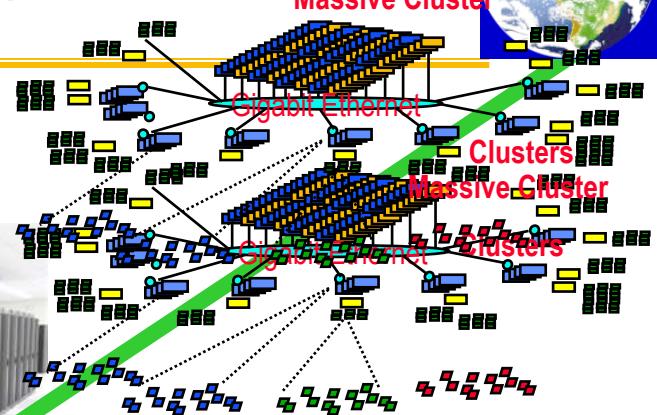


Internet
Connectivity

Scalable, Reliable,
Secure Services

Databases
Information Collection
Remote Storage
Online Games
Commerce
...

MEMS for
Sensor Nets



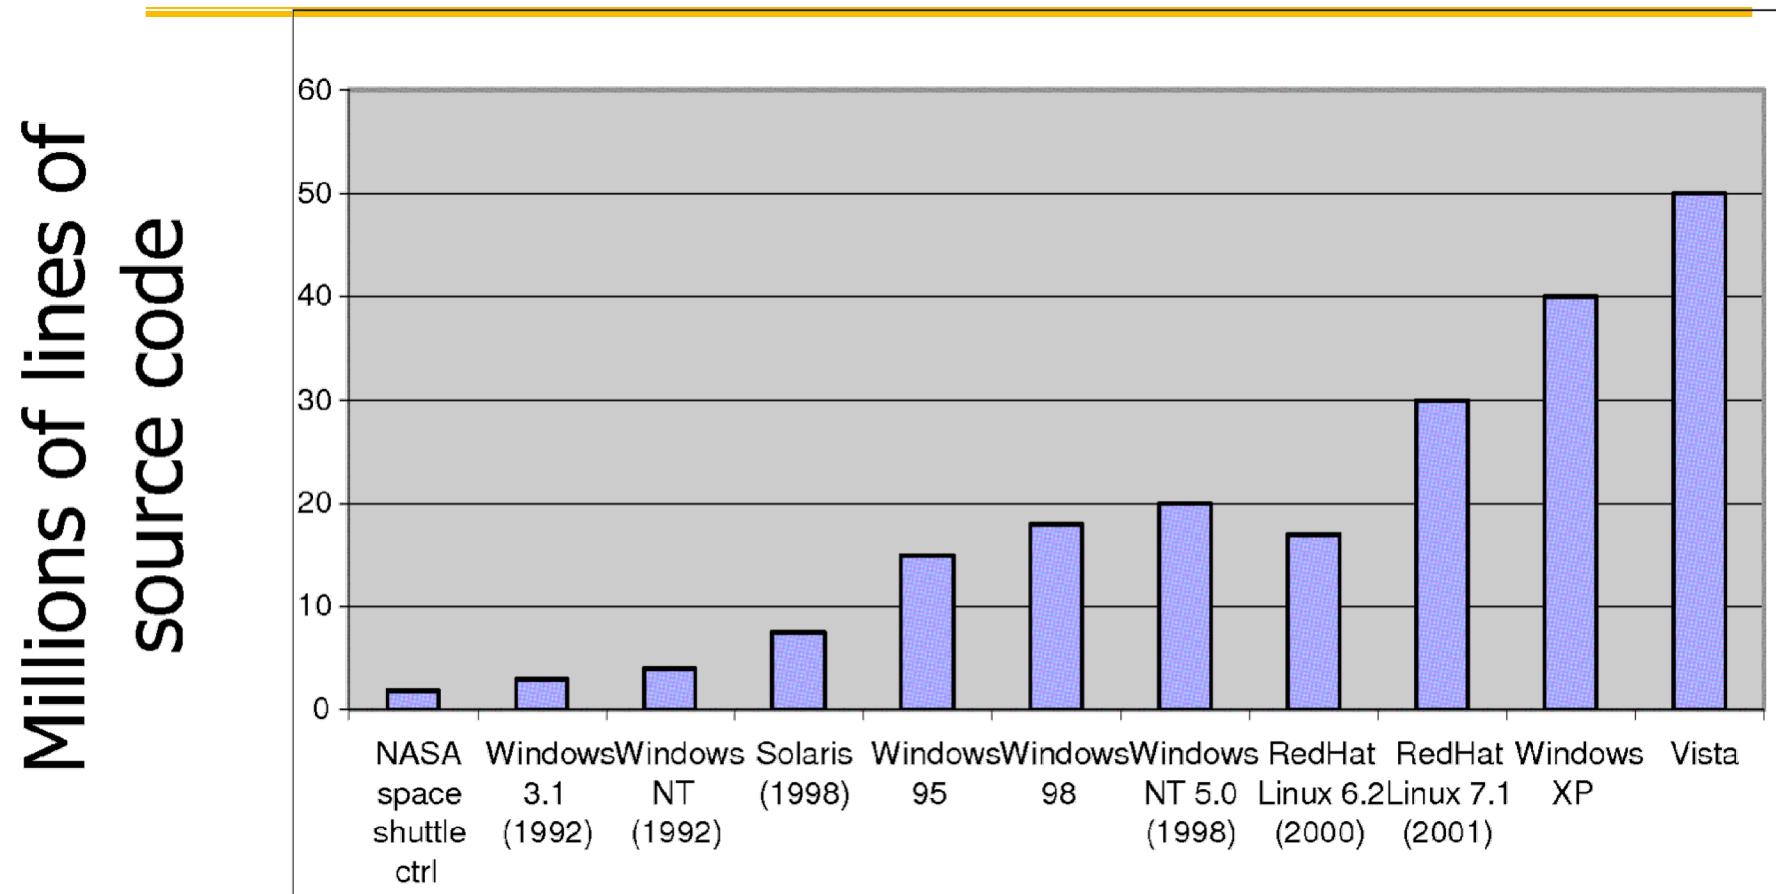


Challenge: Complexity

- **Applications consisting of...**
 - ... a variety of software modules that ...
 - ... run on a variety of devices (machines) that
 - » ... implement different hardware architectures
 - » ... run competing applications
 - » ... fail in unexpected ways
 - » ... can be under a variety of attacks
- **Not feasible to test software for all possible environments and combinations of components and devices**
 - The question is not whether there are bugs but how well are they managed.



Increasing Software Complexity



From MIT's 6.033 course



How do We Tame Complexity?

- Every piece of computer hardware different
 - Different CPU
 - » Pentium, ARM, PowerPC, ColdFire
 - Different amounts of memory, disk, ...
 - Different types of devices
 - » Mice, keyboards, sensors, cameras, fingerprint readers, touch screen
 - Different networking environment
 - » Cable, DSL, Wireless, ...
- Questions:
 - Does the programmer need to write a single program that performs many independent activities?
 - Does every program have to be altered for every piece of hardware?
 - Does a faulty program crash everything?

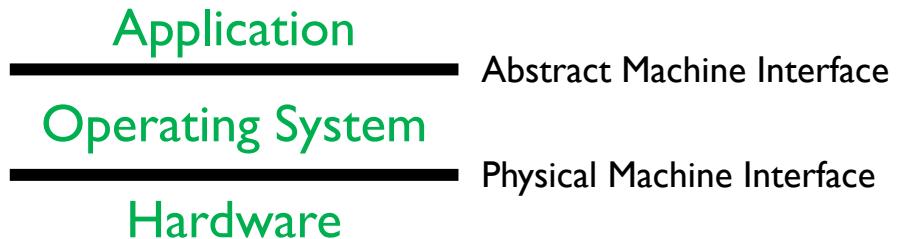


Virtualization ??



OS *Abstracts* underlying hardware

- Processor => Thread
- Memory => Address Space
- Disks, SSDs, ... => Files
- Networks => Sockets
- Machines => Processes

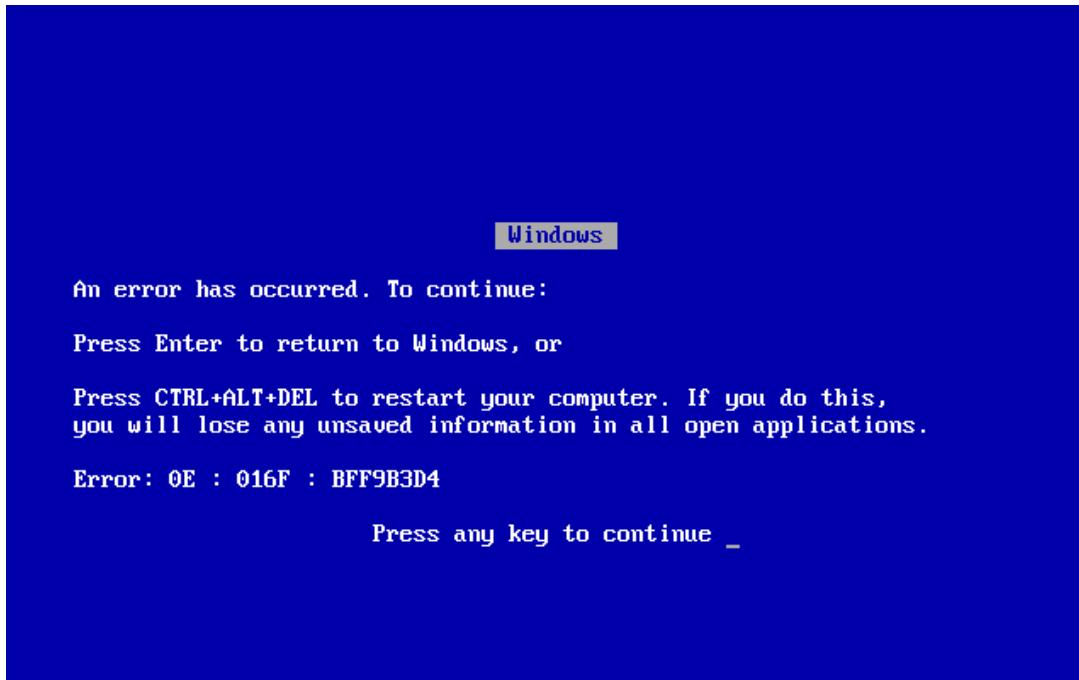


- OS Goals:
 - Remove software/hardware quirks (*fight complexity*)
 - Optimize for convenience, utilization, reliability, ... (*help the programmer*)
- For any OS area (e.g. file systems, virtual memory, networking, scheduling):
 - What hardware interface to handle? (physical reality)
 - What's software interface to provide? (nicer abstraction)



OS Goal: Protecting Processes & The Kernel

- Run multiple applications and:
 - Keep them from interfering with or crashing the operating system
 - Keep them from interfering with or crashing each other





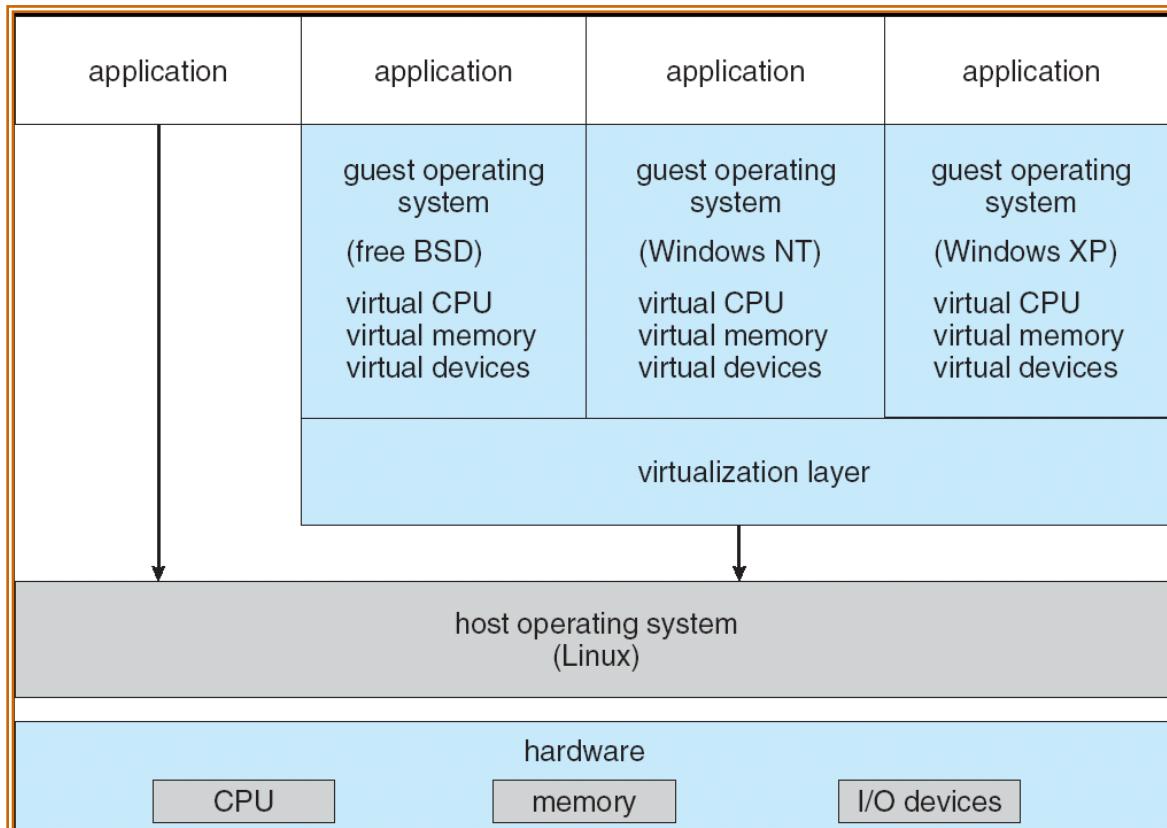
Virtual Machines

- Virtualize every detail of a hardware configuration so perfectly that you can run an operating system (and many applications) on top of it.
 - VMWare Fusion, Virtual box, Parallels Desktop, Xen, Vagrant
- Provides isolation
- Complete insulation from change
- The norm in the Cloud (server consolidation)
- Long history (60's in IBM OS development)
- All our work will take place INSIDE a VM
 - Vagrant (new image just for you)

System Virtual Machines: Layers of OSs

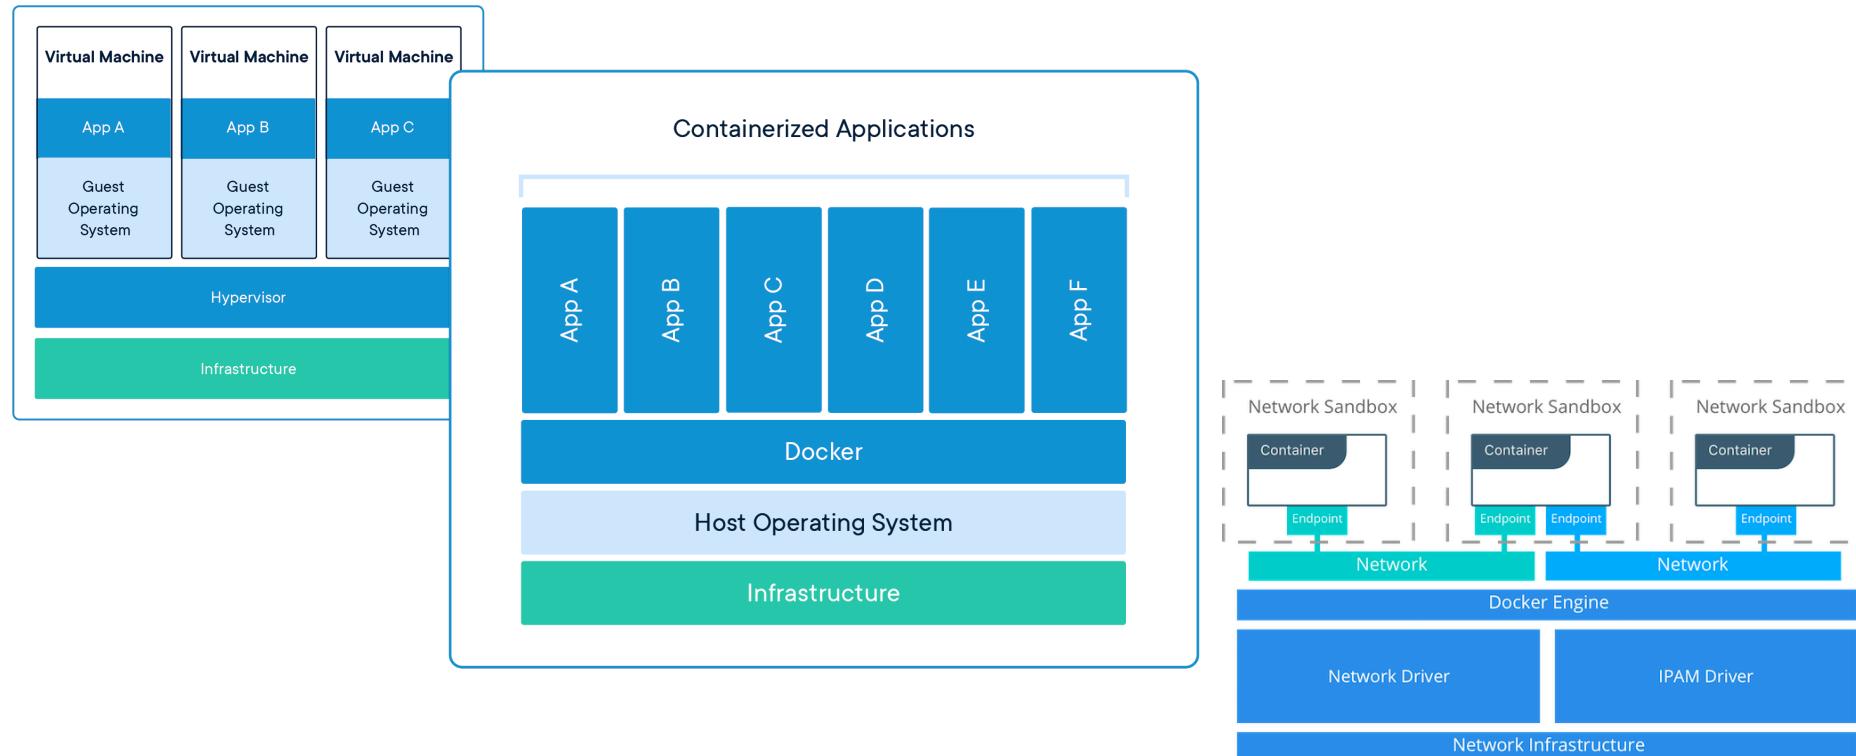


- Useful for OS development
 - When OS crashes, restricted to one VM
 - Can aid testing programs on other OSs





Containers *virtualize the OS*

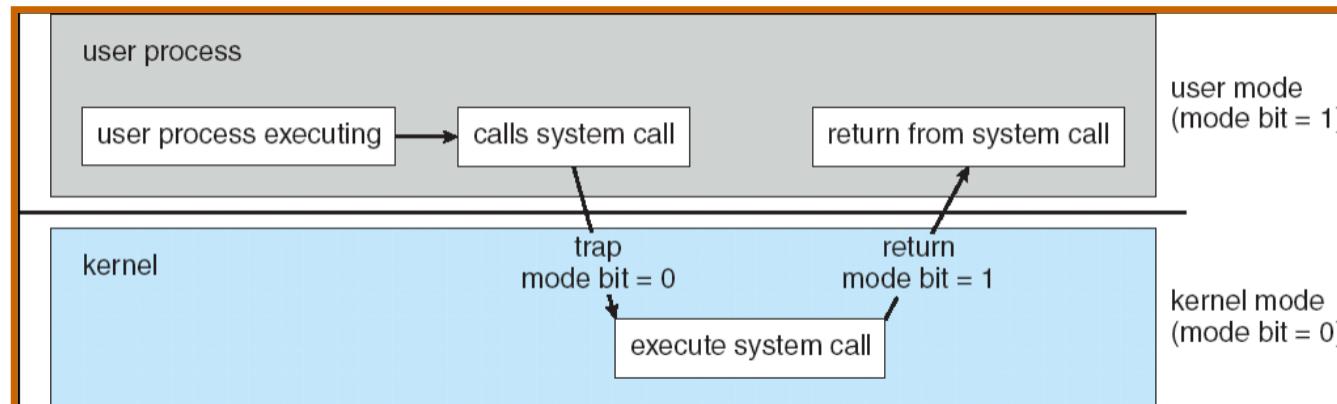


- Roots in OS developments to provide protected systems abstraction, not just application abstraction
 - User-level file system (route syscalls to user process)
 - Cgroups – predictable, bounded resources (CPU, Mem, BW)



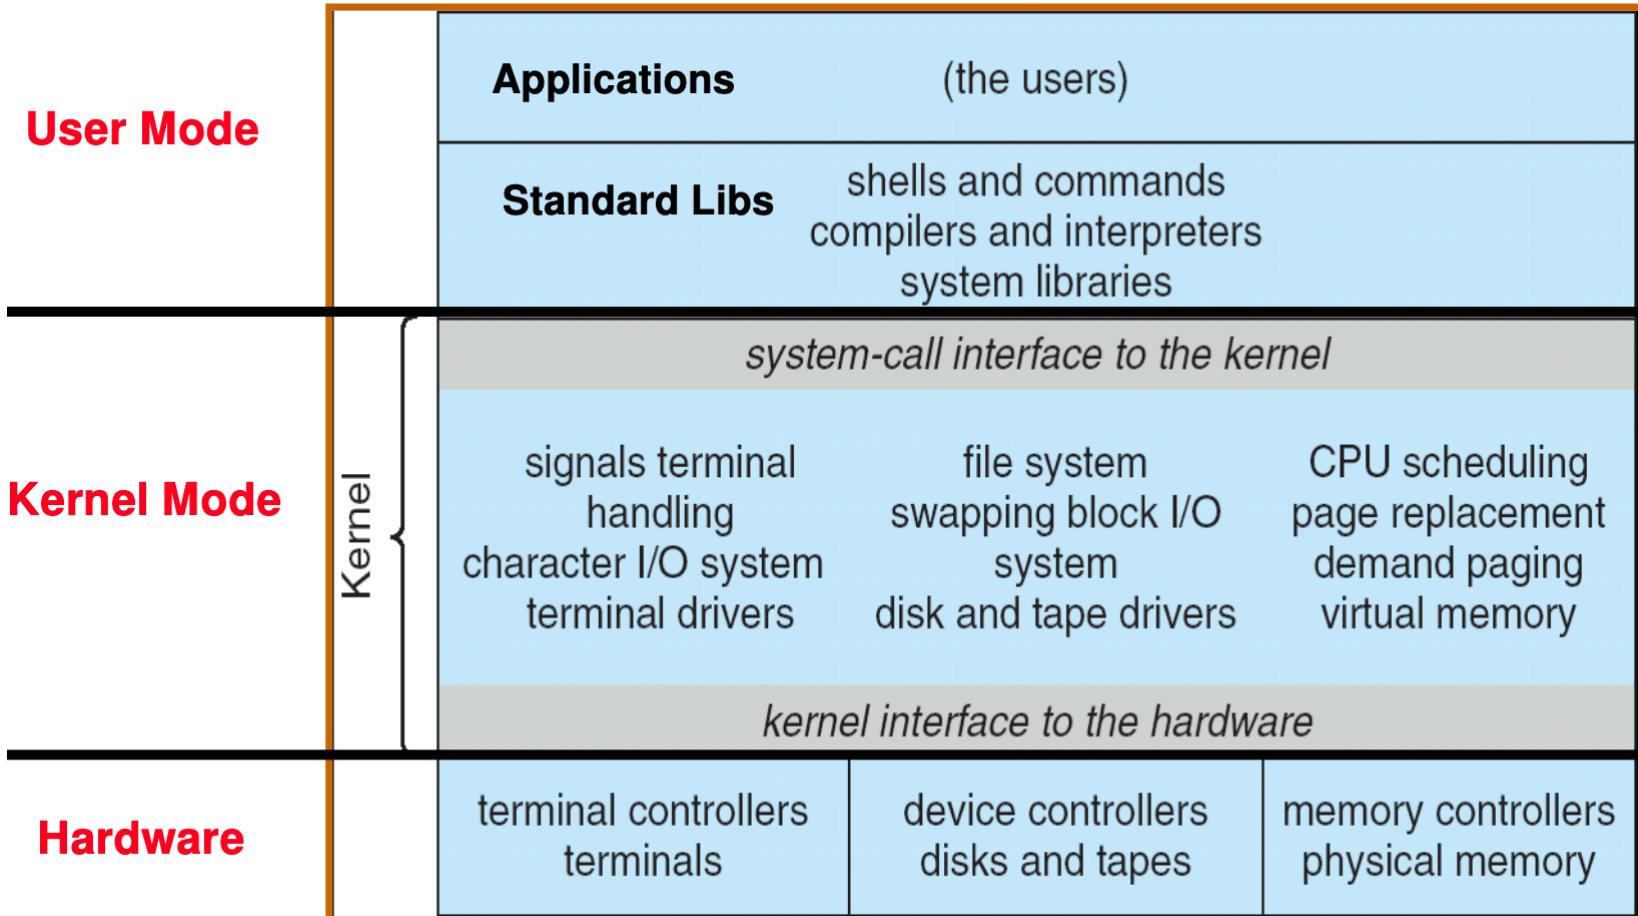
Basic tool: Dual Mode Operation

- **Hardware provides at least two modes:**
 1. Kernel Mode (or "supervisor" / "protected" mode)
 2. User Mode
- **Certain operations are prohibited when running in user mode**
 - Changing the page table pointer
- **Carefully controlled transitions between user mode and kernel mode**
 - System calls, interrupts, exceptions





UNIX OS Structure





Summary: Operating Systems...

- **Provide a machine abstraction to handle diverse hardware**
 - Convenience, protection, reliability obtain in creating the illusion
- **Coordinate resources and protect users from each other**
 - Utilizing a few critical hardware mechanisms
- **Simplify application development by providing standard services**
- **Provide fault containment, fault tolerance, and fault recovery**
- **Key concepts: concurrency, isolation, protected sharing, virtualization**